

NAME

adjtimex, clock_adjtime, ntp_adjtime – tune kernel clock

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <sys/timex.h>
```

```
int adjtimex(struct timex *buf);
```

```
int clock_adjtime(clockid_t clk_id, struct timex *buf);
```

```
int ntp_adjtime(struct timex *buf);
```

DESCRIPTION

Linux uses David L. Mills' clock adjustment algorithm (see RFC 5905). The system call **adjtimex()** reads and optionally sets adjustment parameters for this algorithm. It takes a pointer to a *timex* structure, updates kernel parameters from (selected) field values, and returns the same structure updated with the current kernel values. This structure is declared as follows:

```
struct timex {
    int  modes;           /* Mode selector */
    long offset;          /* Time offset; nanoseconds, if STA_NANO
                           status flag is set, otherwise
                           microseconds */
    long freq;            /* Frequency offset; see NOTES for units */
    long maxerror;        /* Maximum error (microseconds) */
    long esterror;        /* Estimated error (microseconds) */
    int  status;          /* Clock command/status */
    long constant;        /* PLL (phase-locked loop) time constant */
    long precision;       /* Clock precision
                           (microseconds, read-only) */
    long tolerance;       /* Clock frequency tolerance (read-only);
                           see NOTES for units */
    struct timeval time;  /* Current time (read-only, except for
                           ADJ_SETOFFSET); upon return, time.tv_usec
                           contains nanoseconds, if STA_NANO status
                           flag is set, otherwise microseconds */
    long tick;            /* Microseconds between clock ticks */
    long ppsfreq;         /* PPS (pulse per second) frequency
                           (read-only); see NOTES for units */
    long jitter;          /* PPS jitter (read-only); nanoseconds, if
                           STA_NANO status flag is set, otherwise
                           microseconds */
    int  shift;           /* PPS interval duration
                           (seconds, read-only) */
    long stabil;          /* PPS stability (read-only);
                           see NOTES for units */
    long jitcnt;          /* PPS count of jitter limit exceeded
                           events (read-only) */
    long calcnt;          /* PPS count of calibration intervals
                           (read-only) */
    long errcnt;          /* PPS count of calibration errors
                           (read-only) */
    long stbcnt;          /* PPS count of stability limit exceeded
                           events (read-only) */
    int  tai;             /* TAI offset, as set by previous ADJ_TAI
```

```

        operation (seconds, read-only,
        since Linux 2.6.26) */
    /* Further padding bytes to allow for future expansion */
};

```

The *modes* field determines which parameters, if any, to set. (As described later in this page, the constants used for **ntp_adjtime()** are equivalent but differently named.) It is a bit mask containing a bitwise-*or* combination of zero or more of the following bits:

ADJ_OFFSET

Set time offset from *buf.offset*. Since Linux 2.6.26, the supplied value is clamped to the range (−0.5s, +0.5s). In older kernels, an **EINVAL** error occurs if the supplied value is out of range.

ADJ_FREQUENCY

Set frequency offset from *buf.freq*. Since Linux 2.6.26, the supplied value is clamped to the range (−32768000, +32768000). In older kernels, an **EINVAL** error occurs if the supplied value is out of range.

ADJ_MAXERROR

Set maximum time error from *buf.maxerror*.

ADJ_ESTERROR

Set estimated time error from *buf.esterror*.

ADJ_STATUS

Set clock status bits from *buf.status*. A description of these bits is provided below.

ADJ_TIMECONST

Set PLL time constant from *buf.constant*. If the **STA_NANO** status flag (see below) is clear, the kernel adds 4 to this value.

ADJ_SETOFFSET (since Linux 2.6.39)

Add *buf.time* to the current time. If *buf.status* includes the **ADJ_NANO** flag, then *buf.time.tv_usec* is interpreted as a nanosecond value; otherwise it is interpreted as microseconds.

The value of *buf.time* is the sum of its two fields, but the field *buf.time.tv_usec* must always be nonnegative. The following example shows how to normalize a *timeval* with nanosecond resolution.

```

while (buf.time.tv_usec < 0) {
    buf.time.tv_sec  -= 1;
    buf.time.tv_usec += 1000000000;
}

```

ADJ_MICRO (since Linux 2.6.26)

Select microsecond resolution.

ADJ_NANO (since Linux 2.6.26)

Select nanosecond resolution. Only one of **ADJ_MICRO** and **ADJ_NANO** should be specified.

ADJ_TAI (since Linux 2.6.26)

Set TAI (Atomic International Time) offset from *buf.constant*.

ADJ_TAI should not be used in conjunction with **ADJ_TIMECONST**, since the latter mode also employs the *buf.constant* field.

For a complete explanation of TAI and the difference between TAI and UTC, see *BIPM* (<http://www.bipm.org/en/bipm/tai/tai.html>)

ADJ_TICK

Set tick value from *buf.tick*.

Alternatively, *modes* can be specified as either of the following (multibit mask) values, in which case other bits should not be specified in *modes*:

ADJ_OFFSET_SINGLES

Old-fashioned **adjtime(3)**: (gradually) adjust time by value specified in *buf.offset*, which specifies an adjustment in microseconds.

ADJ_OFFSET_SS_READ (functional since Linux 2.6.28)

Return (in *buf.offset*) the remaining amount of time to be adjusted after an earlier **ADJ_OFFSET_SINGLES** operation. This feature was added in Linux 2.6.24, but did not work correctly until Linux 2.6.28.

Ordinary users are restricted to a value of either 0 or **ADJ_OFFSET_SS_READ** for *modes*. Only the superuser may set any parameters.

The *buf.status* field is a bit mask that is used to set and/or retrieve status bits associated with the NTP implementation. Some bits in the mask are both readable and settable, while others are read-only.

STA_PLL (read-write)

Enable phase-locked loop (PLL) updates via **ADJ_OFFSET**.

STA_PPSFREQ (read-write)

Enable PPS (pulse-per-second) frequency discipline.

STA_PPSTIME (read-write)

Enable PPS time discipline.

STA_FLL (read-write)

Select frequency-locked loop (FLL) mode.

STA_INS (read-write)

Insert a leap second after the last second of the UTC day, thus extending the last minute of the day by one second. Leap-second insertion will occur each day, so long as this flag remains set.

STA_DEL (read-write)

Delete a leap second at the last second of the UTC day. Leap second deletion will occur each day, so long as this flag remains set.

STA_UNSYNC (read-write)

Clock unsynchronized.

STA_FREQHOLD (read-write)

Hold frequency. Normally adjustments made via **ADJ_OFFSET** result in dampened frequency adjustments also being made. So a single call corrects the current offset, but as offsets in the same direction are made repeatedly, the small frequency adjustments will accumulate to fix the long-term skew.

This flag prevents the small frequency adjustment from being made when correcting for an **ADJ_OFFSET** value.

STA_PPSSIGNAL (read-only)

A valid PPS (pulse-per-second) signal is present.

STA_PPSJITTER (read-only)

PPS signal jitter exceeded.

STA_PPSWANDER (read-only)

PPS signal wander exceeded.

STA_PPSERROR (read-only)

PPS signal calibration error.

STA_CLOCKERR (read-only)

Clock hardware fault.

STA_NANO (read-only; since Linux 2.6.26)

Resolution (0 = microsecond, 1 = nanoseconds). Set via **ADJ_NANO**, cleared via **ADJ_MICRO**.

STA_MODE (since Linux 2.6.26)

Mode (0 = Phase Locked Loop, 1 = Frequency Locked Loop).

STA_CLK (read-only; since Linux 2.6.26)

Clock source (0 = A, 1 = B); currently unused.

Attempts to set read-only *status* bits are silently ignored.

clock_adjtime ()

The **clock_adjtime()** system call (added in Linux 2.6.39) behaves like **adjtimex()** but takes an additional *clk_id* argument to specify the particular clock on which to act.

ntp_adjtime ()

The **ntp_adjtime()** library function (described in the NTP "Kernel Application Program API", KAPI) is a more portable interface for performing the same task as **adjtimex()**. Other than the following points, it is identical to **adjtimex()**:

- The constants used in *modes* are prefixed with "MOD_" rather than "ADJ_", and have the same suffixes (thus, **MOD_OFFSET**, **MOD_FREQUENCY**, and so on), other than the exceptions noted in the following points.
- **MOD_CLKA** is the synonym for **ADJ_OFFSET_SINGLESHOT**.
- **MOD_CLKB** is the synonym for **ADJ_TICK**.
- There is no synonym for **ADJ_OFFSET_SS_READ**, which is not described in the KAPI.

RETURN VALUE

On success, **adjtimex()** and **ntp_adjtime()** return the clock state; that is, one of the following values:

TIME_OK Clock synchronized, no leap second adjustment pending.

TIME_INS Indicates that a leap second will be added at the end of the UTC day.

TIME_DEL Indicates that a leap second will be deleted at the end of the UTC day.

TIME_OOP Insertion of a leap second is in progress.

TIME_WAIT

A leap-second insertion or deletion has been completed. This value will be returned until the next **ADJ_STATUS** operation clears the **STA_INS** and **STA_DEL** flags.

TIME_ERROR

The system clock is not synchronized to a reliable server. This value is returned when any of the following holds true:

- Either **STA_UNSYNC** or **STA_CLOCKERR** is set.
- **STA_PPSSIGNAL** is clear and either **STA_PPSFREQ** or **STA_PPSTIME** is set.
- **STA_PPSTIME** and **STA_PPSJITTER** are both set.
- **STA_PPSFREQ** is set and either **STA_PPSWANDER** or **STA_PPSJITTER** is set.

The symbolic name **TIME_BAD** is a synonym for **TIME_ERROR**, provided for backward compatibility.

Note that starting with Linux 3.4, the call operates asynchronously and the return value usually will not reflect a state change caused by the call itself.

On failure, these calls return -1 and set *errno* to indicate the error.

ERRORS

EFAULT

buf does not point to writable memory.

EINVAL (before Linux 2.6.26)

An attempt was made to set *buf.freq* to a value outside the range $(-33554432, +33554432)$.

EINVAL (before Linux 2.6.26)

An attempt was made to set *buf.offset* to a value outside the permitted range. Before Linux 2.0, the permitted range was (-131072, +131072). From Linux 2.0 onwards, the permitted range was (-512000, +512000).

EINVAL

An attempt was made to set *buf.status* to a value other than those listed above.

EINVAL

The *clk_id* given to **clock_adjtime()** is invalid for one of two reasons. Either the System-V style hard-coded positive clock ID value is out of range, or the dynamic *clk_id* does not refer to a valid instance of a clock object. See **clock_gettime(2)** for a discussion of dynamic clocks.

EINVAL

An attempt was made to set *buf.tick* to a value outside the range 900000/**HZ** to 1100000/**HZ**, where **HZ** is the system timer interrupt frequency.

ENODEV

The hot-pluggable device (like USB for example) represented by a dynamic *clk_id* has disappeared after its character device was opened. See **clock_gettime(2)** for a discussion of dynamic clocks.

EOPNOTSUPP

The given *clk_id* does not support adjustment.

EPERM

buf.modes is neither 0 nor **ADJ_OFFSET_SS_READ**, and the caller does not have sufficient privilege. Under Linux, the **CAP_SYS_TIME** capability is required.

ATTRIBUTES

For an explanation of the terms used in this section, see **attributes(7)**.

Interface	Attribute	Value
ntp_adjtime()	Thread safety	MT-Safe

STANDARDS

None of these interfaces is described in POSIX.1

adjtimex() and **clock_adjtime()** are Linux-specific and should not be used in programs intended to be portable.

The preferred API for the NTP daemon is **ntp_adjtime()**.

NOTES

In struct *timex*, *freq*, *ppsfreq*, and *stabil* are ppm (parts per million) with a 16-bit fractional part, which means that a value of 1 in one of those fields actually means 2^{-16} ppm, and $2^{16}=65536$ is 1 ppm. This is the case for both input values (in the case of *freq*) and output values.

The leap-second processing triggered by **STA_INS** and **STA_DEL** is done by the kernel in timer context. Thus, it will take one tick into the second for the leap second to be inserted or deleted.

SEE ALSO

clock_gettime(2), **clock_settime(2)**, **settimeofday(2)**, **adjtime(3)**, **ntp_gettime(3)**, **capabilities(7)**, **time(7)**, **adjtimex(8)**, **hwclock(8)**

NTP "Kernel Application Program Interface" (<http://www.slac.stanford.edu/comp/unix/package/rtems/src/ssrApps/ntpNanoclock/api.htm>)