

NAME

arch_prctl – set architecture-specific thread state

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <asm/prctl.h>    /* Definition of ARCH_* constants */
#include <sys/syscall.h>   /* Definition of SYS_* constants */
#include <unistd.h>
```

```
int syscall(SYS_arch_prctl, int code, unsigned long addr);
int syscall(SYS_arch_prctl, int code, unsigned long *addr);
```

Note: glibc provides no wrapper for **arch_prctl()**, necessitating the use of **syscall(2)**.

DESCRIPTION

arch_prctl() sets architecture-specific process or thread state. *code* selects a subfunction and passes argument *addr* to it; *addr* is interpreted as either an *unsigned long* for the "set" operations, or as an *unsigned long **, for the "get" operations.

Subfunctions for both x86 and x86-64 are:

ARCH_SET_CPUID (since Linux 4.12)

Enable (*addr != 0*) or disable (*addr == 0*) the *cpuid* instruction for the calling thread. The instruction is enabled by default. If disabled, any execution of a *cpuid* instruction will instead generate a **SIGSEGV** signal. This feature can be used to emulate *cpuid* results that differ from what the underlying hardware would have produced (e.g., in a paravirtualization setting).

The **ARCH_SET_CPUID** setting is preserved across **fork(2)** and **clone(2)** but reset to the default (i.e., *cpuid* enabled) on **execve(2)**.

ARCH_GET_CPUID (since Linux 4.12)

Return the setting of the flag manipulated by **ARCH_SET_CPUID** as the result of the system call (1 for enabled, 0 for disabled). *addr* is ignored.

Subfunctions for x86-64 only are:

ARCH_SET_FS

Set the 64-bit base for the *FS* register to *addr*.

ARCH_GET_FS

Return the 64-bit base value for the *FS* register of the calling thread in the *unsigned long* pointed to by *addr*.

ARCH_SET_GS

Set the 64-bit base for the *GS* register to *addr*.

ARCH_GET_GS

Return the 64-bit base value for the *GS* register of the calling thread in the *unsigned long* pointed to by *addr*.

RETURN VALUE

On success, **arch_prctl()** returns 0; on error, *-1* is returned, and *errno* is set to indicate the error.

ERRORS**EFAULT**

addr points to an unmapped address or is outside the process address space.

EINVAL

code is not a valid subcommand.

ENODEV

ARCH_SET_CPUID was requested, but the underlying hardware does not support CPUID faulting.

EPERM

addr is outside the process address space.

STANDARDS

arch_prctl() is a Linux/x86-64 extension and should not be used in programs intended to be portable.

NOTES

arch_prctl() is supported only on Linux/x86-64 for 64-bit programs currently.

The 64-bit base changes when a new 32-bit segment selector is loaded.

ARCH_SET_GS is disabled in some kernels.

Context switches for 64-bit segment bases are rather expensive. As an optimization, if a 32-bit TLS base address is used, **arch_prctl()** may use a real TLS entry as if **set_thread_area(2)** had been called, instead of manipulating the segment base register directly. Memory in the first 2 GB of address space can be allocated by using **mmap(2)** with the **MAP_32BIT** flag.

Because of the aforementioned optimization, using **arch_prctl()** and **set_thread_area(2)** in the same thread is dangerous, as they may overwrite each other's TLS entries.

FS may be already used by the threading library. Programs that use **ARCH_SET_FS** directly are very likely to crash.

SEE ALSO

mmap(2), **modify_ldt(2)**, **prctl(2)**, **set_thread_area(2)**

AMD X86-64 Programmer's manual