

**NAME**

btrfs-qgroup – control the quota group of a btrfs filesystem

**SYNOPSIS**

**btrfs qgroup** <subcommand> <args>

**DESCRIPTION**

**btrfs qgroup** is used to control quota group (qgroup) of a btrfs filesystem.

**Note**

To use qgroup you need to enable quota first using **btrfs quota enable** command.

**Warning**

Qgroup is not stable yet and will impact performance in current mainline kernel (v4.14).

**QGROUP**

Quota groups or qgroup in btrfs make a tree hierarchy, the leaf qgroups are attached to subvolumes. The size limits are set per qgroup and apply when any limit is reached in tree that contains a given subvolume.

The limits are separated between shared and exclusive and reflect the extent ownership. For example a fresh snapshot shares almost all the blocks with the original subvolume, new writes to either subvolume will raise towards the exclusive limit.

The qgroup identifiers conform to *level/id* where level 0 is reserved to the qgroups associated with subvolumes. Such qgroups are created automatically.

The qgroup hierarchy is built by commands **create** and **assign**.

**Note**

If the qgroup of a subvolume is destroyed, quota about the subvolume will not be functional until qgroup 0/<subvolume id> is created again.

**SUBCOMMAND**

**assign** [options] <src> <dst> <path>

Assign qgroup <src> as the child qgroup of <dst> in the btrfs filesystem identified by <path>.

**Options**

--rescan

(default since: 4.19) Automatically schedule quota rescan if the new qgroup assignment would lead to quota inconsistency. See *QUOTA RESCAN* for more information.

--no-rescan

Explicitly ask not to do a rescan, even if the assignment will make the quotas inconsistent. This may be useful for repeated calls where the rescan would add unnecessary overhead.

**create** <qgroupid> <path>

Create a subvolume quota group.

For the 0/<subvolume id> qgroup, a qgroup can be created even before the subvolume is created.

**destroy** <qgroupid> <path>

Destroy a qgroup.

If a qgroup is not isolated, meaning it is a parent or child qgroup, then it can only be destroyed after the relationship is removed.

**limit** [options] <size>|none [<qgroupid>] <path>

Limit the size of a qgroup to <size> or no limit in the btrfs filesystem identified by <path>.

If <qgroupid> is not given, qgroup of the subvolume identified by <path> is used if possible.

**Options**

- c  
limit amount of data after compression. This is the default, it is currently not possible to turn off this option.
- e  
limit space exclusively assigned to this qgroup.

**remove** <src> <dst> <path>

Remove the relationship between child qgroup <src> and parent qgroup <dst> in the btrfs filesystem identified by <path>.

**Options**

- rescan  
(default since: 4.19) Automatically schedule quota rescan if the removed qgroup relation would lead to quota inconsistency. See *QUOTA RESCAN* for more information.
- no-rescan  
Explicitly ask not to do a rescan, even if the removal will make the quotas inconsistent. This may be useful for repeated calls where the rescan would add unnecessary overhead.

**show** [options] <path>

Show all qgroups in the btrfs filesystem identified by <path>.

**Options**

- p  
print parent qgroup id.
- c  
print child qgroup id.
- r  
print limit of referenced size of qgroup.
- e  
print limit of exclusive size of qgroup.
- F  
list all qgroups which impact the given path(include ancestral qgroups)
- f  
list all qgroups which impact the given path(exclude ancestral qgroups)
- raw  
raw numbers in bytes, without the *B* suffix.
- human-readable  
print human friendly numbers, base 1024, this is the default
- iec  
select the 1024 base for the following options, according to the IEC standard.
- si  
select the 1000 base for the following options, according to the SI standard.
- kbytes  
show sizes in KiB, or kB with --si.
- mbytes  
show sizes in MiB, or MB with --si.
- gbytes  
show sizes in GiB, or GB with --si.

`--bytes`  
show sizes in TiB, or TB with `--si`.

`--sort=[+/-]<attr>[, [+/-]<attr>]...`  
list qgroups in order of `<attr>`.

`<attr>` can be one or more of `qgroupid`, `rfer`, `excl`, `max_rfer`, `max_excl`.

Prefix '+' means ascending order and '-' means descending order of `<attr>`. If no prefix is given, use ascending order by default.

If multiple `<attr>`s is given, use comma to separate.

`--sync`  
To retrieve information after updating the state of qgroups, force sync of the filesystem identified by `<path>` before getting information.

## QUOTA RESCAN

The rescan reads all extent sharing metadata and updates the respective qgroups accordingly.

The information consists of bytes owned exclusively (*excl*) or shared/referred to (*rfer*). There's no explicit information about which extents are shared or owned exclusively. This means when qgroup relationship changes, extent owners change and qgroup numbers are no longer consistent unless we do a full rescan.

However there are cases where we can avoid a full rescan, if a subvolume whose *rfer* number equals its *excl* number, which means all bytes are exclusively owned, then assigning/removing this subvolume only needs to add/subtract *rfer* number from its parent qgroup. This can speed up the rescan.

## EXAMPLES

### Example 1. Make a parent group that has two quota group children

Given the following filesystem mounted at `/mnt/my-vault`

```
Label: none  uuid: 60d2ab3b-941a-4f22-8d1a-315f329797b2
Total devices 1 FS bytes used 128.00KiB
devid    1 size 5.00GiB used 536.00MiB path /dev/vdb
```

Enable quota and create subvolumes. Check subvolume ids.

```
$ cd /mnt/my-vault
$ btrfs quota enable .
$ btrfs subvolume create a
$ btrfs subvolume create b
$ btrfs subvolume list .
```

```
ID 261 gen 61 top level 5 path a
ID 262 gen 62 top level 5 path b
```

Create qgroup and set limit to 10MiB.

```
$ btrfs qgroup create 1/100 .
$ btrfs qgroup limit 10M 1/100 .
$ btrfs qgroup assign 0/261 1/100 .
$ btrfs qgroup assign 0/262 1/100 .
```

And check qgroups.

```
$ btrfs qgroup show .
```

qgroupid	rfer	excl
-----	-----	-----
0/5	16.00KiB	16.00KiB
0/261	16.00KiB	16.00KiB
0/262	16.00KiB	16.00KiB
1/100	32.00KiB	32.00KiB

## EXIT STATUS

**btrfs qgroup** returns a zero exit status if it succeeds. Non zero is returned in case of failure.

## AVAILABILITY

**btrfs** is part of **btrfs-progs**. Please refer to the btrfs wiki <http://btrfs.wiki.kernel.org> for further details.

## SEE ALSO

**mkfs.btrfs(8)**, **btrfs-subvolume(8)**, **btrfs-quota(8)**,