

NAME

fanotify_mark – add, remove, or modify an fanotify mark on a filesystem object

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <sys/fanotify.h>
```

```
int fanotify_mark(int fanotify_fd, unsigned int flags,
                  uint64_t mask, int dirfd,
                  const char *_Nullable pathname);
```

DESCRIPTION

For an overview of the fanotify API, see **fanotify(7)**.

fanotify_mark() adds, removes, or modifies an fanotify mark on a filesystem object. The caller must have read permission on the filesystem object that is to be marked.

The *fanotify_fd* argument is a file descriptor returned by **fanotify_init(2)**.

flags is a bit mask describing the modification to perform. It must include exactly one of the following values:

FAN_MARK_ADD

The events in *mask* will be added to the mark mask (or to the ignore mask). *mask* must be nonempty or the error **EINVAL** will occur.

FAN_MARK_REMOVE

The events in argument *mask* will be removed from the mark mask (or from the ignore mask). *mask* must be nonempty or the error **EINVAL** will occur.

FAN_MARK_FLUSH

Remove either all marks for filesystems, all marks for mounts, or all marks for directories and files from the fanotify group. If *flags* contains **FAN_MARK_MOUNT**, all marks for mounts are removed from the group. If *flags* contains **FAN_MARK_FILESYSTEM**, all marks for filesystems are removed from the group. Otherwise, all marks for directories and files are removed. No flag other than, and at most one of, the flags **FAN_MARK_MOUNT** or **FAN_MARK_FILESYSTEM** can be used in conjunction with **FAN_MARK_FLUSH**. *mask* is ignored.

If none of the values above is specified, or more than one is specified, the call fails with the error **EINVAL**.

In addition, zero or more of the following values may be ORed into *flags*:

FAN_MARK_DONT_FOLLOW

If *pathname* is a symbolic link, mark the link itself, rather than the file to which it refers. (By default, **fanotify_mark()** dereferences *pathname* if it is a symbolic link.)

FAN_MARK_ONLYDIR

If the filesystem object to be marked is not a directory, the error **ENOTDIR** shall be raised.

FAN_MARK_MOUNT

Mark the mount specified by *pathname*. If *pathname* is not itself a mount point, the mount containing *pathname* will be marked. All directories, subdirectories, and the contained files of the mount will be monitored. The events which require that filesystem objects are identified by file handles, such as **FAN_CREATE**, **FAN_ATTRIB**, **FAN_MOVE**, and **FAN_DELETE_SELF**, cannot be provided as a *mask* when *flags* contains **FAN_MARK_MOUNT**. Attempting to do so will result in the error **EINVAL** being returned. Use of this flag requires the **CAP_SYS_ADMIN** capability.

FAN_MARK_FILESYSTEM (since Linux 4.20)

Mark the filesystem specified by *pathname*. The filesystem containing *pathname* will be marked. All the contained files and directories of the filesystem from any mount point will be monitored. Use of this flag requires the **CAP_SYS_ADMIN** capability.

FAN_MARK_IGNORED_MASK

The events in *mask* shall be added to or removed from the ignore mask. Note that the flags **FAN_ONDIR**, and **FAN_EVENT_ON_CHILD** have no effect when provided with this flag. The effect of setting the flags **FAN_ONDIR**, and **FAN_EVENT_ON_CHILD** in the mark mask on the events that are set in the ignore mask is undefined and depends on the Linux kernel version. Specifically, prior to Linux 5.9, setting a mark mask on a file and a mark with ignore mask on its parent directory would not result in ignoring events on the file, regardless of the **FAN_EVENT_ON_CHILD** flag in the parent directory's mark mask. When the ignore mask is updated with the **FAN_MARK_IGNORED_MASK** flag on a mark that was previously updated with the **FAN_MARK_IGNORE** flag, the update fails with **EEXIST** error.

FAN_MARK_IGNORE (since Linux 6.0)

This flag has a similar effect as setting the **FAN_MARK_IGNORED_MASK** flag. The events in *mask* shall be added to or removed from the ignore mask. Unlike the **FAN_MARK_IGNORED_MASK** flag, this flag also has the effect that the **FAN_ONDIR**, and **FAN_EVENT_ON_CHILD** flags take effect on the ignore mask. Specifically, unless the **FAN_ONDIR** flag is set with **FAN_MARK_IGNORE**, events on directories will not be ignored. If the flag **FAN_EVENT_ON_CHILD** is set with **FAN_MARK_IGNORE**, events on children will be ignored. For example, a mark on a directory with combination of a mask with **FAN_CREATE** event and **FAN_ONDIR** flag and an ignore mask with **FAN_CREATE** event and without **FAN_ONDIR** flag, will result in getting only the events for creation of sub-directories. When using the **FAN_MARK_IGNORE** flag to add to an ignore mask of a mount, filesystem, or directory inode mark, the **FAN_MARK_IGNORED_SURV_MODIFY** flag must be specified. Failure to do so will result with **EINVAL** or **EISDIR** error.

FAN_MARK_IGNORED_SURV_MODIFY

The ignore mask shall survive modify events. If this flag is not set, the ignore mask is cleared when a modify event occurs on the marked object. Omitting this flag is typically used to suppress events (e.g., **FAN_OPEN**) for a specific file, until that specific file's content has been modified. It is far less useful to suppress events on an entire filesystem, or mount, or on all files inside a directory, until some file's content has been modified. For this reason, the **FAN_MARK_IGNORE** flag requires the **FAN_MARK_IGNORED_SURV_MODIFY** flag on a mount, filesystem, or directory inode mark. This flag cannot be removed from a mark once set. When the ignore mask is updated without this flag on a mark that was previously updated with the **FAN_MARK_IGNORE** and **FAN_MARK_IGNORED_SURV_MODIFY** flags, the update fails with **EEXIST** error.

FAN_MARK_IGNORE_SURV

This is a synonym for (**FAN_MARK_IGNORE**|**FAN_MARK_IGNORED_SURV_MODIFY**).

FAN_MARK_EVICTABLE (since Linux 5.19)

When an inode mark is created with this flag, the inode object will not be pinned to the inode cache, therefore, allowing the inode object to be evicted from the inode cache when the memory pressure on the system is high. The eviction of the inode object results in the evictable mark also being lost. When the mask of an evictable inode mark is updated without using the **FAN_MARK_EVICTABLE** flag, the marked inode is pinned to inode cache and the mark is no longer evictable. When the mask of a non-evictable inode mark is updated with the **FAN_MARK_EVICTABLE** flag, the inode mark remains non-evictable and the update fails with **EEXIST** error. Mounts and filesystems are not evictable objects, therefore, an attempt to create a mount mark or a filesystem mark with the **FAN_MARK_EVICTABLE** flag, will result in the error **EINVAL**. For example, inode marks can be used in combination with mount marks to reduce the amount of events from noninteresting paths. The event listener reads events, checks if the path reported in the event is of interest, and if it is not, the listener sets a mark with an ignore mask on the directory. Evictable inode marks allow using this method for a large number of directories without the concern of pinning all inodes and exhausting the system's memory.

mask defines which events shall be listened for (or which shall be ignored). It is a bit mask composed of the following values:

FAN_ACCESS

Create an event when a file or directory (but see BUGS) is accessed (read).

FAN_MODIFY

Create an event when a file is modified (write).

FAN_CLOSE_WRITE

Create an event when a writable file is closed.

FAN_CLOSE_NOWRITE

Create an event when a read-only file or directory is closed.

FAN_OPEN

Create an event when a file or directory is opened.

FAN_OPEN_EXEC (since Linux 5.0)

Create an event when a file is opened with the intent to be executed. See NOTES for additional details.

FAN_ATTRIB (since Linux 5.1)

Create an event when the metadata for a file or directory has changed. An fanotify group that identifies filesystem objects by file handles is required.

FAN_CREATE (since Linux 5.1)

Create an event when a file or directory has been created in a marked parent directory. An fanotify group that identifies filesystem objects by file handles is required.

FAN_DELETE (since Linux 5.1)

Create an event when a file or directory has been deleted in a marked parent directory. An fanotify group that identifies filesystem objects by file handles is required.

FAN_DELETE_SELF (since Linux 5.1)

Create an event when a marked file or directory itself is deleted. An fanotify group that identifies filesystem objects by file handles is required.

FAN_FS_ERROR (since Linux 5.16)

Create an event when a filesystem error leading to inconsistent filesystem metadata is detected. An additional information record of type **FAN_EVENT_INFO_TYPE_ERROR** is returned for each event in the read buffer. An fanotify group that identifies filesystem objects by file handles is required.

Events of such type are dependent on support from the underlying filesystem. At the time of writing, only the **ext4** filesystem reports **FAN_FS_ERROR** events.

See **fanotify(7)** for additional details.

FAN_MOVED_FROM (since Linux 5.1)

Create an event when a file or directory has been moved from a marked parent directory. An fanotify group that identifies filesystem objects by file handles is required.

FAN_MOVED_TO (since Linux 5.1)

Create an event when a file or directory has been moved to a marked parent directory. An fanotify group that identifies filesystem objects by file handles is required.

FAN_RENAME (since Linux 5.17)

This event contains the same information provided by events **FAN_MOVED_FROM** and **FAN_MOVED_TO**, however is represented by a single event with up to two information records. An fanotify group that identifies filesystem objects by file handles is required. If the filesystem object to be marked is not a directory, the error **ENOTDIR** shall be raised.

FAN_MOVE_SELF (since Linux 5.1)

Create an event when a marked file or directory itself has been moved. An fanotify group that identifies filesystem objects by file handles is required.

FAN_OPEN_PERM

Create an event when a permission to open a file or directory is requested. An fanotify file descriptor created with **FAN_CLASS_PRE_CONTENT** or **FAN_CLASS_CONTENT** is required.

FAN_OPEN_EXEC_PERM (since Linux 5.0)

Create an event when a permission to open a file for execution is requested. An fanotify file descriptor created with **FAN_CLASS_PRE_CONTENT** or **FAN_CLASS_CONTENT** is required. See NOTES for additional details.

FAN_ACCESS_PERM

Create an event when a permission to read a file or directory is requested. An fanotify file descriptor created with **FAN_CLASS_PRE_CONTENT** or **FAN_CLASS_CONTENT** is required.

FAN_ONDIR

Create events for directories—for example, when **opendir(3)**, **readdir(3)** (but see BUGS), and **closedir(3)** are called. Without this flag, events are created only for files. In the context of directory entry events, such as **FAN_CREATE**, **FAN_DELETE**, **FAN_MOVED_FROM**, and **FAN_MOVED_TO**, specifying the flag **FAN_ONDIR** is required in order to create events when subdirectory entries are modified (i.e., **mkdir(2)**/**rmdir(2)**).

FAN_EVENT_ON_CHILD

Events for the immediate children of marked directories shall be created. The flag has no effect when marking mounts and filesystems. Note that events are not generated for children of the subdirectories of marked directories. More specifically, the directory entry modification events **FAN_CREATE**, **FAN_DELETE**, **FAN_MOVED_FROM**, and **FAN_MOVED_TO** are not generated for any entry modifications performed inside subdirectories of marked directories. Note that the events **FAN_DELETE_SELF** and **FAN_MOVE_SELF** are not generated for children of marked directories. To monitor complete directory trees it is necessary to mark the relevant mount or filesystem.

The following composed values are defined:

FAN_CLOSE

A file is closed (**FAN_CLOSE_WRITE**|**FAN_CLOSE_NOWRITE**).

FAN_MOVE

A file or directory has been moved (**FAN_MOVED_FROM**|**FAN_MOVED_TO**).

The filesystem object to be marked is determined by the file descriptor *dirfd* and the pathname specified in *pathname*:

- If *pathname* is NULL, *dirfd* defines the filesystem object to be marked.
- If *pathname* is NULL, and *dirfd* takes the special value **AT_FDCWD**, the current working directory is to be marked.
- If *pathname* is absolute, it defines the filesystem object to be marked, and *dirfd* is ignored.
- If *pathname* is relative, and *dirfd* does not have the value **AT_FDCWD**, then the filesystem object to be marked is determined by interpreting *pathname* relative to the directory referred to by *dirfd*.
- If *pathname* is relative, and *dirfd* has the value **AT_FDCWD**, then the filesystem object to be marked is determined by interpreting *pathname* relative to the current working directory. (See **openat(2)** for an explanation of why the *dirfd* argument is useful.)

RETURN VALUE

On success, **fanotify_mark()** returns 0. On error, **-1** is returned, and *errno* is set to indicate the error.

ERRORS**EBADF**

An invalid file descriptor was passed in *fanotify_fd*.

EBADF

pathname is relative but *dirfd* is neither **AT_FDCWD** nor a valid file descriptor.

EEXIST

The filesystem object indicated by *dirfd* and *pathname* has a mark that was updated without the **FAN_MARK_EVICTABLE** flag, and the user attempted to update the mark with **FAN_MARK_EVICTABLE** flag.

EEXIST

The filesystem object indicated by *dirfd* and *pathname* has a mark that was updated with the **FAN_MARK_IGNORE** flag, and the user attempted to update the mark with **FAN_MARK_IGNORED_MASK** flag.

EEXIST

The filesystem object indicated by *dirfd* and *pathname* has a mark that was updated with the **FAN_MARK_IGNORE** and **FAN_MARK_IGNORED_SURV_MODIFY** flags, and the user attempted to update the mark only with **FAN_MARK_IGNORE** flag.

EINVAL

An invalid value was passed in *flags* or *mask*, or *fanotify_fd* was not an fanotify file descriptor.

EINVAL

The fanotify file descriptor was opened with **FAN_CLASS_NOTIF** or the fanotify group identifies filesystem objects by file handles and *mask* contains a flag for permission events (**FAN_OPEN_PERM** or **FAN_ACCESS_PERM**).

EINVAL

The group was initialized without **FAN_REPORT_FID** but one or more event types specified in the *mask* require it.

EINVAL

flags contains **FAN_MARK_IGNORE**, and either **FAN_MARK_MOUNT** or **FAN_MARK_FILESYSTEM**, but does not contain **FAN_MARK_IGNORED_SURV_MODIFY**.

EISDIR

flags contains **FAN_MARK_IGNORE**, but does not contain **FAN_MARK_IGNORED_SURV_MODIFY**, and *dirfd* and *pathname* specify a directory.

ENODEV

The filesystem object indicated by *dirfd* and *pathname* is not associated with a filesystem that supports *fsid* (e.g., **fuse(4)**). **tmpfs(5)** did not support *fsid* prior to Linux 5.13. This error can be returned only with an fanotify group that identifies filesystem objects by file handles.

ENOENT

The filesystem object indicated by *dirfd* and *pathname* does not exist. This error also occurs when trying to remove a mark from an object which is not marked.

ENOMEM

The necessary memory could not be allocated.

ENOSPC

The number of marks for this user exceeds the limit and the **FAN_UNLIMITED_MARKS** flag was not specified when the fanotify file descriptor was created with **fanotify_init(2)**. See **fanotify(7)** for details about this limit.

ENOSYS

This kernel does not implement **fanotify_mark()**. The fanotify API is available only if the kernel was configured with **CONFIG_FANOTIFY**.

ENOTDIR

flags contains **FAN_MARK_ONLYDIR**, and *dirfd* and *pathname* do not specify a directory.

ENOTDIR

mask contains **FAN_RENAME**, and *dirfd* and *pathname* do not specify a directory.

ENOTDIR

flags contains **FAN_MARK_IGNORE**, or the fanotify group was initialized with flag **FAN_REPORT_TARGET_FID**, and *mask* contains directory entry modification events (e.g., **FAN_CREATE**, **FAN_DELETE**), or directory event flags (e.g., **FAN_ONDIR**, **FAN_EVENT_ON_CHILD**), and *dirfd* and *pathname* do not specify a directory.

EOPNOTSUPP

The object indicated by *pathname* is associated with a filesystem that does not support the encoding of file handles. This error can be returned only with an fanotify group that identifies filesystem objects by file handles.

EPERM

The operation is not permitted because the caller lacks a required capability.

EXDEV

The filesystem object indicated by *pathname* resides within a filesystem sub volume (e.g., **btrfs**(5)) which uses a different *fsid* than its root superblock. This error can be returned only with an fanotify group that identifies filesystem objects by file handles.

VERSIONS

fanotify_mark() was introduced in Linux 2.6.36 and enabled in Linux 2.6.37.

STANDARDS

This system call is Linux-specific.

NOTES**FAN_OPEN_EXEC and FAN_OPEN_EXEC_PERM**

When using either **FAN_OPEN_EXEC** or **FAN_OPEN_EXEC_PERM** within the *mask*, events of these types will be returned only when the direct execution of a program occurs. More specifically, this means that events of these types will be generated for files that are opened using **execve**(2), **execveat**(2), or **uselib**(2). Events of these types will not be raised in the situation where an interpreter is passed (or reads) a file for interpretation.

Additionally, if a mark has also been placed on the Linux dynamic linker, a user should also expect to receive an event for it when an ELF object has been successfully opened using **execve**(2) or **execveat**(2).

For example, if the following ELF binary were to be invoked and a **FAN_OPEN_EXEC** mark has been placed on /:

```
$ /bin/echo foo
```

The listening application in this case would receive **FAN_OPEN_EXEC** events for both the ELF binary and interpreter, respectively:

```
/bin/echo
/lib64/ld-linux-x86-64.so.2
```

BUGS

The following bugs were present in before Linux 3.16:

- If *flags* contains **FAN_MARK_FLUSH**, *dirfd*, and *pathname* must specify a valid filesystem object, even though this object is not used.
- **readdir**(2) does not generate a **FAN_ACCESS** event.
- If **fanotify_mark**() is called with **FAN_MARK_FLUSH**, *flags* is not checked for invalid values.

SEE ALSO

fanotify_init(2), **fanotify**(7)