

**NAME**

flock – apply or remove an advisory lock on an open file

**LIBRARY**

Standard C library (*libc*, *-lc*)

**SYNOPSIS**

```
#include <sys/file.h>
```

```
int flock(int fd, int operation);
```

**DESCRIPTION**

Apply or remove an advisory lock on the open file specified by *fd*. The argument *operation* is one of the following:

**LOCK\_SH**

Place a shared lock. More than one process may hold a shared lock for a given file at a given time.

**LOCK\_EX**

Place an exclusive lock. Only one process may hold an exclusive lock for a given file at a given time.

**LOCK\_UN**

Remove an existing lock held by this process.

A call to **flock()** may block if an incompatible lock is held by another process. To make a nonblocking request, include **LOCK\_NB** (by ORing) with any of the above operations.

A single file may not simultaneously have both shared and exclusive locks.

Locks created by **flock()** are associated with an open file description (see **open(2)**). This means that duplicate file descriptors (created by, for example, **fork(2)** or **dup(2)**) refer to the same lock, and this lock may be modified or released using any of these file descriptors. Furthermore, the lock is released either by an explicit **LOCK\_UN** operation on any of these duplicate file descriptors, or when all such file descriptors have been closed.

If a process uses **open(2)** (or similar) to obtain more than one file descriptor for the same file, these file descriptors are treated independently by **flock()**. An attempt to lock the file using one of these file descriptors may be denied by a lock that the calling process has already placed via another file descriptor.

A process may hold only one type of lock (shared or exclusive) on a file. Subsequent **flock()** calls on an already locked file will convert an existing lock to the new lock mode.

Locks created by **flock()** are preserved across an **execve(2)**.

A shared or exclusive lock can be placed on a file regardless of the mode in which the file was opened.

**RETURN VALUE**

On success, zero is returned. On error, *-1* is returned, and *errno* is set to indicate the error.

**ERRORS****EBADF**

*fd* is not an open file descriptor.

**EINTR**

While waiting to acquire a lock, the call was interrupted by delivery of a signal caught by a handler; see **signal(7)**.

**EINVAL**

*operation* is invalid.

**ENOLCK**

The kernel ran out of memory for allocating lock records.

**EWOULDBLOCK**

The file is locked and the **LOCK\_NB** flag was selected.

**STANDARDS**

4.4BSD (the **flock()** call first appeared in 4.2BSD). A version of **flock()**, possibly implemented in terms of **fcntl(2)**, appears on most UNIX systems.

**NOTES**

Since Linux 2.0, **flock()** is implemented as a system call in its own right rather than being emulated in the GNU C library as a call to **fcntl(2)**. With this implementation, there is no interaction between the types of lock placed by **flock()** and **fcntl(2)**, and **flock()** does not detect deadlock. (Note, however, that on some systems, such as the modern BSDs, **flock()** and **fcntl(2)** locks *do* interact with one another.)

**flock()** places advisory locks only; given suitable permissions on a file, a process is free to ignore the use of **flock()** and perform I/O on the file.

**flock()** and **fcntl(2)** locks have different semantics with respect to forked processes and **dup(2)**. On systems that implement **flock()** using **fcntl(2)**, the semantics of **flock()** will be different from those described in this manual page.

Converting a lock (shared to exclusive, or vice versa) is not guaranteed to be atomic: the existing lock is first removed, and then a new lock is established. Between these two steps, a pending lock request by another process may be granted, with the result that the conversion either blocks, or fails if **LOCK\_NB** was specified. (This is the original BSD behavior, and occurs on many other implementations.)

**NFS details**

Up to Linux 2.6.11, **flock()** does not lock files over NFS (i.e., the scope of locks was limited to the local system). Instead, one could use **fcntl(2)** byte-range locking, which does work over NFS, given a sufficiently recent version of Linux and a server which supports locking.

Since Linux 2.6.12, NFS clients support **flock()** locks by emulating them as **fcntl(2)** byte-range locks on the entire file. This means that **fcntl(2)** and **flock()** locks *do* interact with one another over NFS. It also means that in order to place an exclusive lock, the file must be opened for writing.

Since Linux 2.6.37, the kernel supports a compatibility mode that allows **flock()** locks (and also **fcntl(2)** byte region locks) to be treated as local; see the discussion of the *local\_lock* option in **nfs(5)**.

**CIFS details**

Up to Linux 5.4, **flock()** is not propagated over SMB. A file with such locks will not appear locked for remote clients.

Since Linux 5.5, **flock()** locks are emulated with SMB byte-range locks on the entire file. Similarly to NFS, this means that **fcntl(2)** and **flock()** locks interact with one another. Another important side-effect is that the locks are not advisory anymore: any IO on a locked file will always fail with **EACCES** when done from a separate file descriptor. This difference originates from the design of locks in the SMB protocol, which provides mandatory locking semantics.

Remote and mandatory locking semantics may vary with SMB protocol, mount options and server type. See **mount.cifs(8)** for additional information.

**SEE ALSO**

**flock(1)**, **close(2)**, **dup(2)**, **execve(2)**, **fcntl(2)**, **fork(2)**, **open(2)**, **lockf(3)**, **lslocks(8)**

*Documentation/filesystems/locks.txt* in the Linux kernel source tree (*Documentation/locks.txt* in older kernels)