

**NAME**

get\_thread\_area, set\_thread\_area – manipulate thread-local storage information

**LIBRARY**

Standard C library (*libc*, *-lc*)

**SYNOPSIS**

```
#include <sys/syscall.h> /* Definition of SYS_* constants */
#include <unistd.h>

#ifdef __i386__ || defined __x86_64__
# include <asm/ldt.h> /* Definition of struct user_desc */

int syscall(SYS_get_thread_area, struct user_desc *u_info);
int syscall(SYS_set_thread_area, struct user_desc *u_info);

#elif defined __m68k__

int syscall(SYS_get_thread_area);
int syscall(SYS_set_thread_area, unsigned long tp);

#elif defined __mips__

int syscall(SYS_set_thread_area, unsigned long addr);

#endif
```

*Note:* glibc provides no wrappers for these system calls, necessitating the use of **syscall(2)**.

**DESCRIPTION**

These calls provide architecture-specific support for a thread-local storage implementation. At the moment, **set\_thread\_area()** is available on m68k, MIPS, and x86 (both 32-bit and 64-bit variants); **get\_thread\_area()** is available on m68k and x86.

On m68k and MIPS, **set\_thread\_area()** allows storing an arbitrary pointer (provided in the **tp** argument on m68k and in the **addr** argument on MIPS) in the kernel data structure associated with the calling thread; this pointer can later be retrieved using **get\_thread\_area()** (see also NOTES for information regarding obtaining the thread pointer on MIPS).

On x86, Linux dedicates three global descriptor table (GDT) entries for thread-local storage. For more information about the GDT, see the Intel Software Developer's Manual or the AMD Architecture Programming Manual.

Both of these system calls take an argument that is a pointer to a structure of the following type:

```
struct user_desc {
    unsigned int  entry_number;
    unsigned int  base_addr;
    unsigned int  limit;
    unsigned int  seg_32bit:1;
    unsigned int  contents:2;
    unsigned int  read_exec_only:1;
    unsigned int  limit_in_pages:1;
    unsigned int  seg_not_present:1;
    unsigned int  useable:1;
#ifdef __x86_64__
    unsigned int  lm:1;
#endif
};
```

**get\_thread\_area()** reads the GDT entry indicated by *u\_info->entry\_number* and fills in the rest of the fields in *u\_info*.

**set\_thread\_area()** sets a TLS entry in the GDT.

The TLS array entry set by **set\_thread\_area()** corresponds to the value of *u\_info->entry\_number* passed in by the user. If this value is in bounds, **set\_thread\_area()** writes the TLS descriptor pointed to by *u\_info* into the thread's TLS array.

When **set\_thread\_area()** is passed an *entry\_number* of `-1`, it searches for a free TLS entry. If **set\_thread\_area()** finds a free TLS entry, the value of *u\_info->entry\_number* is set upon return to show which entry was changed.

A *user\_desc* is considered "empty" if *read\_exec\_only* and *seg\_not\_present* are set to 1 and all of the other fields are 0. If an "empty" descriptor is passed to **set\_thread\_area()**, the corresponding TLS entry will be cleared. See BUGS for additional details.

Since Linux 3.19, **set\_thread\_area()** cannot be used to write non-present segments, 16-bit segments, or code segments, although clearing a segment is still acceptable.

## RETURN VALUE

On x86, these system calls return 0 on success, and `-1` on failure, with *errno* set to indicate the error.

On MIPS and m68k, **set\_thread\_area()** always returns 0. On m68k, **get\_thread\_area()** returns the thread area pointer value (previously set via **set\_thread\_area()**).

## ERRORS

### EFAULT

*u\_info* is an invalid pointer.

### EINVAL

*u\_info->entry\_number* is out of bounds.

### ENOSYS

**get\_thread\_area()** or **set\_thread\_area()** was invoked as a 64-bit system call.

### ESRCH

(**set\_thread\_area()**) A free TLS entry could not be located.

## VERSIONS

**set\_thread\_area()** first appeared in Linux 2.5.29. **get\_thread\_area()** first appeared in Linux 2.5.32.

## STANDARDS

**set\_thread\_area()** and **get\_thread\_area()** are Linux-specific and should not be used in programs that are intended to be portable.

## NOTES

These system calls are generally intended for use only by threading libraries.

**arch\_prctl(2)** can interfere with **set\_thread\_area()** on x86. See **arch\_prctl(2)** for more details. This is not normally a problem, as **arch\_prctl(2)** is normally used only by 64-bit programs.

On MIPS, the current value of the thread area pointer can be obtained using the instruction:

```
rdhwr dest, $29
```

This instruction traps and is handled by kernel.

## BUGS

On 64-bit kernels before Linux 3.19, one of the padding bits in *user\_desc*, if set, would prevent the descriptor from being considered empty (see **modify\_ldt(2)**). As a result, the only reliable way to clear a TLS entry is to use **memset(3)** to zero the entire *user\_desc* structure, including padding bits, and then to set the *read\_exec\_only* and *seg\_not\_present* bits. On Linux 3.19, a *user\_desc* consisting entirely of zeros except for *entry\_number* will also be interpreted as a request to clear a TLS entry, but this behaved differently on older kernels.

Prior to Linux 3.19, the DS and ES segment registers must not reference TLS entries.

## SEE ALSO

**arch\_prctl(2)**, **modify\_ldt(2)**, **ptrace(2)** (**PTRACE\_GET\_THREAD\_AREA** and **PTRACE\_SET\_THREAD\_AREA**)