

**NAME**

getdents, getdents64 – get directory entries

**LIBRARY**

Standard C library (*libc*, *-lc*)

**SYNOPSIS**

```
#include <sys/syscall.h>    /* Definition of SYS_* constants */
#include <unistd.h>

long syscall(SYS_getdents, unsigned int fd, struct linux_dirent *dirp,
             unsigned int count);

#define _GNU_SOURCE          /* See feature_test_macros(7) */
#include <dirent.h>

ssize_t getdents64(int fd, void dirp[count], size_t count);
```

*Note:* glibc provides no wrapper for **getdents()**, necessitating the use of **syscall(2)**.

*Note:* There is no definition of *struct linux\_dirent* in glibc; see NOTES.

**DESCRIPTION**

These are not the interfaces you are interested in. Look at **readdir(3)** for the POSIX-conforming C library interface. This page documents the bare kernel system call interfaces.

**getdents()**

The system call **getdents()** reads several *linux\_dirent* structures from the directory referred to by the open file descriptor *fd* into the buffer pointed to by *dirp*. The argument *count* specifies the size of that buffer.

The *linux\_dirent* structure is declared as follows:

```
struct linux_dirent {
    unsigned long d_ino;        /* Inode number */
    unsigned long d_off;        /* Offset to next linux_dirent */
    unsigned short d_reclen;    /* Length of this linux_dirent */
    char          d_name[];     /* Filename (null-terminated) */
                                /* length is actually (d_reclen - 2 -
                                /*      offsetof(struct linux_dirent, d_name)) */
    /*
    char          pad;          // Zero padding byte
    char          d_type;       // File type (only since Linux
                                // 2.6.4); offset is (d_reclen - 1)
    */
}
```

*d\_ino* is an inode number. *d\_off* is the distance from the start of the directory to the start of the next *linux\_dirent*. *d\_reclen* is the size of this entire *linux\_dirent*. *d\_name* is a null-terminated filename.

*d\_type* is a byte at the end of the structure that indicates the file type. It contains one of the following values (defined in *<dirent.h>*):

<b>DT_BLK</b>	This is a block device.
<b>DT_CHR</b>	This is a character device.
<b>DT_DIR</b>	This is a directory.
<b>DT_FIFO</b>	This is a named pipe (FIFO).
<b>DT_LNK</b>	This is a symbolic link.
<b>DT_REG</b>	This is a regular file.
<b>DT SOCK</b>	This is a UNIX domain socket.

**DT\_UNKNOWN**

The file type is unknown.

The *d\_type* field is implemented since Linux 2.6.4. It occupies a space that was previously a zero-filled padding byte in the *linux\_dirent* structure. Thus, on kernels up to and including Linux 2.6.3, attempting to access this field always provides the value 0 (**DT\_UNKNOWN**).

Currently, only some filesystems (among them: Btrfs, ext2, ext3, and ext4) have full support for returning the file type in *d\_type*. All applications must properly handle a return of **DT\_UNKNOWN**.

**getdents64()**

The original Linux **getdents()** system call did not handle large filesystems and large file offsets. Consequently, Linux 2.4 added **getdents64()**, with wider types for the *d\_ino* and *d\_off* fields. In addition, **getdents64()** supports an explicit *d\_type* field.

The **getdents64()** system call is like **getdents()**, except that its second argument is a pointer to a buffer containing structures of the following type:

```
struct linux_dirent64 {
    ino64_t      d_ino;      /* 64-bit inode number */
    off64_t      d_off;      /* 64-bit offset to next structure */
    unsigned short d_reclen; /* Size of this dirent */
    unsigned char d_type;    /* File type */
    char         d_name[]; /* Filename (null-terminated) */
};
```

**RETURN VALUE**

On success, the number of bytes read is returned. On end of directory, 0 is returned. On error, -1 is returned, and *errno* is set to indicate the error.

**ERRORS****EBADF**

Invalid file descriptor *fd*.

**EFAULT**

Argument points outside the calling process's address space.

**EINVAL**

Result buffer is too small.

**ENOENT**

No such directory.

**ENOTDIR**

File descriptor does not refer to a directory.

**STANDARDS**

SVr4.

**NOTES**

Library support for **getdents64()** was added in glibc 2.30; glibc does not provide a wrapper for **getdents()**; call **getdents()** (or **getdents64()** on earlier glibc versions) using **syscall(2)**. In that case you will need to define the *linux\_dirent* or *linux\_dirent64* structure yourself.

Probably, you want to use **readdir(3)** instead of these system calls.

These calls supersede **readdir(2)**.

**EXAMPLES**

The program below demonstrates the use of **getdents()**. The following output shows an example of what we see when running this program on an ext2 directory:

```
$ ./a.out /testfs/
----- nread=120 -----
inode#   file type  d_reclen  d_off    d_name
```

2	directory	16	12	.
2	directory	16	24	..
11	directory	24	44	lost+found
12	regular	16	56	a
228929	directory	16	68	sub
16353	directory	16	80	sub2
130817	directory	16	4096	sub3

### Program source

```
#define _GNU_SOURCE
#include <dirent.h>      /* Defines DT_* constants */
#include <err.h>
#include <fcntl.h>
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/syscall.h>
#include <unistd.h>

struct linux_dirent {
    unsigned long    d_ino;
    off_t            d_off;
    unsigned short   d_reclen;
    char             d_name[];
};

#define BUF_SIZE 1024

int
main(int argc, char *argv[])
{
    int            fd;
    char           d_type;
    char           buf[BUF_SIZE];
    long           nread;
    struct linux_dirent *d;

    fd = open(argc > 1 ? argv[1] : ".", O_RDONLY | O_DIRECTORY);
    if (fd == -1)
        err(EXIT_FAILURE, "open");

    for (;;) {
        nread = syscall(SYS_getdents, fd, buf, BUF_SIZE);
        if (nread == -1)
            err(EXIT_FAILURE, "getdents");

        if (nread == 0)
            break;

        printf("----- nread=%ld -----\\n", nread);
        printf("inode#    file type  d_reclen  d_off    d_name\\n");
        for (size_t bpos = 0; bpos < nread;) {
            d = (struct linux_dirent *) (buf + bpos);
            printf("%8lu  ", d->d_ino);
```

```
        d_type = *(buf + bpos + d->d_reclen - 1);
        printf("%-10s ", (d_type == DT_REG) ? "regular" :
                (d_type == DT_DIR) ? "directory" :
                (d_type == DT_FIFO) ? "FIFO" :
                (d_type == DT_SOCK) ? "socket" :
                (d_type == DT_LNK) ? "symlink" :
                (d_type == DT_BLK) ? "block dev" :
                (d_type == DT_CHR) ? "char dev" : "???");
        printf("%4d %10jd  %s\n", d->d_reclen,
                (intmax_t) d->d_off, d->d_name);
        bpos += d->d_reclen;
    }
}

exit(EXIT_SUCCESS);
}
```

**SEE ALSO****readdir(2), readdir(3), inode(7)**