

NAME

io_setup – create an asynchronous I/O context

LIBRARY

Standard C library (*libc*, *-lc*)

Alternatively, Asynchronous I/O library (*libaio*, *-laio*); see NOTES.

SYNOPSIS

```
#include <linux/aio_abi.h>      /* Defines needed types */

long io_setup(unsigned int nr_events, aio_context_t *ctx_idp);
```

Note: There is no glibc wrapper for this system call; see NOTES.

DESCRIPTION

Note: this page describes the raw Linux system call interface. The wrapper function provided by *libaio* uses a different type for the *ctx_idp* argument. See NOTES.

The **io_setup()** system call creates an asynchronous I/O context suitable for concurrently processing *nr_events* operations. The *ctx_idp* argument must not point to an AIO context that already exists, and must be initialized to 0 prior to the call. On successful creation of the AIO context, **ctx_idp* is filled in with the resulting handle.

RETURN VALUE

On success, **io_setup()** returns 0. For the failure return, see NOTES.

ERRORS**EAGAIN**

The specified *nr_events* exceeds the limit of available events, as defined in */proc/sys/fs/aio-max-nr* (see **proc(5)**).

EFAULT

An invalid pointer is passed for *ctx_idp*.

EINVAL

ctx_idp is not initialized, or the specified *nr_events* exceeds internal limits. *nr_events* should be greater than 0.

ENOMEM

Insufficient kernel resources are available.

ENOSYS

io_setup() is not implemented on this architecture.

VERSIONS

The asynchronous I/O system calls first appeared in Linux 2.5.

STANDARDS

io_setup() is Linux-specific and should not be used in programs that are intended to be portable.

NOTES

glibc does not provide a wrapper for this system call. You could invoke it using **syscall(2)**. But instead, you probably want to use the **io_setup()** wrapper function provided by *libaio*.

Note that the *libaio* wrapper function uses a different type (*io_context_t **) for the *ctx_idp* argument. Note also that the *libaio* wrapper does not follow the usual C library conventions for indicating errors: on error it returns a negated error number (the negative of one of the values listed in **ERRORS**). If the system call is invoked via **syscall(2)**, then the return value follows the usual conventions for indicating an error: *-1*, with *errno* set to a (positive) value that indicates the error.

SEE ALSO

io_cancel(2), **io_destroy(2)**, **io_getevents(2)**, **io_submit(2)**, **aio(7)**