

NAME

ioctl – control device

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <sys/ioctl.h>
```

```
int ioctl(int fd, unsigned long request, ...);
```

DESCRIPTION

The **ioctl()** system call manipulates the underlying device parameters of special files. In particular, many operating characteristics of character special files (e.g., terminals) may be controlled with **ioctl()** requests. The argument *fd* must be an open file descriptor.

The second argument is a device-dependent request code. The third argument is an untyped pointer to memory. It's traditionally **char *argp** (from the days before **void *** was valid C), and will be so named for this discussion.

An **ioctl()** *request* has encoded in it whether the argument is an *in* parameter or *out* parameter, and the size of the argument *argp* in bytes. Macros and defines used in specifying an **ioctl()** *request* are located in the file *<sys/ioctl.h>*. See NOTES.

RETURN VALUE

Usually, on success zero is returned. A few **ioctl()** requests use the return value as an output parameter and return a nonnegative value on success. On error, *-1* is returned, and *errno* is set to indicate the error.

ERRORS**EBADF**

fd is not a valid file descriptor.

EFAULT

argp references an inaccessible memory area.

EINVAL

request or *argp* is not valid.

ENOTTY

fd is not associated with a character special device.

ENOTTY

The specified request does not apply to the kind of object that the file descriptor *fd* references.

STANDARDS

No single standard. Arguments, returns, and semantics of **ioctl()** vary according to the device driver in question (the call is used as a catch-all for operations that don't cleanly fit the UNIX stream I/O model).

The **ioctl()** system call appeared in Version 7 AT&T UNIX.

NOTES

In order to use this call, one needs an open file descriptor. Often the **open(2)** call has unwanted side effects, that can be avoided under Linux by giving it the **O_NONBLOCK** flag.

ioctl structure

Ioctl command values are 32-bit constants. In principle these constants are completely arbitrary, but people have tried to build some structure into them.

The old Linux situation was that of mostly 16-bit constants, where the last byte is a serial number, and the preceding byte(s) give a type indicating the driver. Sometimes the major number was used: 0x03 for the **HDIO_*** ioctls, 0x06 for the **LP*** ioctls. And sometimes one or more ASCII letters were used. For example, **TCGETS** has value 0x00005401, with 0x54 = 'T' indicating the terminal driver, and **CYGETTIMEOUT** has value 0x00435906, with 0x43 0x59 = 'C' 'Y' indicating the cyclades driver.

Later (0.98p5) some more information was built into the number. One has 2 direction bits (00: none, 01:

write, 10: read, 11: read/write) followed by 14 size bits (giving the size of the argument), followed by an 8-bit type (collecting the ioctls in groups for a common purpose or a common driver), and an 8-bit serial number.

The macros describing this structure live in `<asm/ioctl.h>` and are `_IO(type,nr)` and `{_IOR, _IOW, _IOWR}(type,nr,size)`. They use `sizeof(size)` so that `size` is a misnomer here: this third argument is a data type.

Note that the size bits are very unreliable: in lots of cases they are wrong, either because of buggy macros using `sizeof(sizeof(struct))`, or because of legacy values.

Thus, it seems that the new structure only gave disadvantages: it does not help in checking, but it causes varying values for the various architectures.

SEE ALSO

execve(2), fcntl(2), ioctl_console(2), ioctl_fat(2), ioctl_ficlone(2), ioctl_ficlonerange(2), ioctl_fid-
eduperange(2), ioctl_fslabel(2), ioctl_getfsmap(2), ioctl_iflags(2), ioctl_ns(2), ioctl_tty(2), ioctl_user-
faultfd(2), open(2), sd(4), tty(4)