

**NAME**

perfmonctl – interface to IA-64 performance monitoring unit

**SYNOPSIS**

```
#include <syscall.h>
#include <perfmon.h>
```

```
long perfmonctl(int fd, int cmd, void arg[.narg], int narg);
```

*Note:* There is no glibc wrapper for this system call; see NOTES.

**DESCRIPTION**

The IA-64-specific **perfmonctl()** system call provides an interface to the PMU (performance monitoring unit). The PMU consists of PMD (performance monitoring data) registers and PMC (performance monitoring control) registers, which gather hardware statistics.

**perfmonctl()** applies the operation *cmd* to the input arguments specified by *arg*. The number of arguments is defined by *narg*. The *fd* argument specifies the perfmon context to operate on.

Supported values for *cmd* are:

**PFM\_CREATE\_CONTEXT**

```
perfmonctl(int fd, PFM_CREATE_CONTEXT, pfarg_context_t *ctxt, 1);
```

Set up a context.

The *fd* parameter is ignored. A new perfmon context is created as specified in *ctxt* and its file descriptor is returned in *ctxt->ctx\_fd*.

The file descriptor can be used in subsequent calls to **perfmonctl()** and can be used to read event notifications (type *pfm\_msg\_t*) using **read(2)**. The file descriptor is pollable using **select(2)**, **poll(2)**, and **epoll(7)**.

The context can be destroyed by calling **close(2)** on the file descriptor.

**PFM\_WRITE\_PMCS**

```
perfmonctl(int fd, PFM_WRITE_PMCS, pfarg_reg_t *pmcs, n);
```

Set PMC registers.

**PFM\_WRITE\_PMDS**

```
perfmonctl(int fd, PFM_WRITE_PMDS, pfarg_reg_t *pmds, n);
```

Set PMD registers.

**PFM\_READ\_PMDS**

```
perfmonctl(int fd, PFM_READ_PMDS, pfarg_reg_t *pmds, n);
```

Read PMD registers.

**PFM\_START**

```
perfmonctl(int fd, PFM_START, NULL, 0);
```

Start monitoring.

**PFM\_STOP**

```
perfmonctl(int fd, PFM_STOP, NULL, 0);
```

Stop monitoring.

**PFM\_LOAD\_CONTEXT**

```
perfmonctl(int fd, PFM_LOAD_CONTEXT, pfarg_load_t *larg, 1);
```

Attach the context to a thread.

**PFM\_UNLOAD\_CONTEXT**

```
perfmonctl(int fd, PFM_UNLOAD_CONTEXT, NULL, 0);
```

Detach the context from a thread.

**PFM\_RESTART**

```
perfmonctl(int fd, PFM_RESTART, NULL, 0);
```

Restart monitoring after receiving an overflow notification.

**PFM\_GET\_FEATURES**

```
perfmonctl(int fd, PFM_GET_FEATURES, pfarg_features_t *arg, 1);
```

**PFM\_DEBUG**

```
perfmonctl(int fd, PFM_DEBUG, val, 0);
```

If *val* is nonzero, enable debugging mode, otherwise disable.

**PFM\_GET\_PMC\_RESET\_VAL**

```
perfmonctl(int fd, PFM_GET_PMC_RESET_VAL, pfarg_reg_t *req, n);
```

Reset PMC registers to default values.

**RETURN VALUE**

**perfmonctl()** returns zero when the operation is successful. On error, `-1` is returned and *errno* is set to indicate the error.

**VERSIONS**

**perfmonctl()** was added in Linux 2.4; it was removed in Linux 5.10.

**STANDARDS**

**perfmonctl()** is Linux-specific and is available only on the IA-64 architecture.

**NOTES**

This system call was broken for many years, and ultimately removed in Linux 5.10.

glibc does not provide a wrapper for this system call; on kernels where it exists, call it using **syscall(2)**.

**SEE ALSO**

**gprof(1)**

The perfmon2 interface specification