

NAME

pppd – Point-to-Point Protocol Daemon

SYNOPSIS

pppd [*options*]

DESCRIPTION

PPP is the protocol used for establishing internet links over dial-up modems, DSL connections, and many other types of point-to-point links. The *pppd* daemon works together with the kernel PPP driver to establish and maintain a PPP link with another system (called the *peer*) and to negotiate Internet Protocol (IP) addresses for each end of the link. Pppd can also authenticate the peer and/or supply authentication information to the peer. PPP can be used with other network protocols besides IP, but such use is becoming increasingly rare.

FREQUENTLY USED OPTIONS

ttyname

Use the serial port called *ttyname* to communicate with the peer. If *ttyname* does not begin with a slash (/), the string "/dev/" is prepended to *ttyname* to form the name of the device to open. If no device name is given, or if the name of the terminal connected to the standard input is given, pppd will use that terminal, and will not fork to put itself in the background. A value for this option from a privileged source cannot be overridden by a non-privileged user.

speed

An option that is a decimal number is taken as the desired baud rate for the serial device. On systems such as 4.4BSD and NetBSD, any speed can be specified. Other systems (e.g. Linux, SunOS) only support the commonly-used baud rates.

asynctest *map*

This option sets the Async-Control-Character-Map (ACCM) for this end of the link. The ACCM is a set of 32 bits, one for each of the ASCII control characters with values from 0 to 31, where a 1 bit indicates that the corresponding control character should not be used in PPP packets sent to this system. The map is encoded as a hexadecimal number (without a leading 0x) where the least significant bit (00000001) represents character 0 and the most significant bit (80000000) represents character 31. Pppd will ask the peer to send these characters as a 2-byte escape sequence. If multiple *asynctest* options are given, the values are ORed together. If *noasynctest* option is given, the default is zero, so pppd will ask the peer not to escape any control characters. To escape transmitted characters, use the *escape* option.

auth

Require the peer to authenticate itself before allowing network packets to be sent or received. This option is the default if the system has a default route. If neither this option nor the *noauth* option is specified, pppd will only allow the peer to use IP addresses to which the system does not already have a route.

call *name*

Read additional options from the file */etc/ppp/peers/name*. This file may contain privileged options, such as *noauth*, even if pppd is not being run by root. The *name* string may not begin with / or include .. as a pathname component. The format of the options file is described below.

connect *script*

Usually there is something which needs to be done to prepare the link before the PPP protocol can be started; for instance, with a dial-up modem, commands need to be sent to the modem to dial the appropriate phone number. This option specifies an command for pppd to execute (by passing it to a shell) before attempting to start PPP negotiation. The *chat* (8) program is often useful here, as it provides a way to send arbitrary strings to a modem and respond to received characters. A value for this option from a privileged source cannot be overridden by a non-privileged user.

crtstcts

Specifies that pppd should set the serial port to use hardware flow control using the RTS and CTS signals in the RS-232 interface. If neither the *crtstcts*, the *nocrtstcts*, the *cdtrcts* nor the *nocdtrcts* option is given, the hardware flow control setting for the serial port is left unchanged. Some serial ports (such as Macintosh serial ports) lack a true RTS output. Such serial ports use this mode to implement unidirectional flow control. The serial port will suspend transmission when requested

by the modem (via CTS) but will be unable to request the modem to stop sending to the computer. This mode retains the ability to use DTR as a modem control line.

defaultroute

Add a default route to the system routing tables, using the peer as the gateway, when IPCP negotiation is successfully completed. This entry is removed when the PPP connection is broken. This option is privileged if the *ndefaultroute* option has been specified.

defaultroute-metric

Define the metric of the *defaultroute* and only add it if there is no other default route with the same metric. With the default value of -1, the route is only added if there is no default route at all.

defaultroute6

Add a default IPv6 route to the system routing tables, using the peer as the gateway, when IPv6CP negotiation is successfully completed. This entry is removed when the PPP connection is broken. This option is privileged if the *ndefaultroute6* option has been specified.

replacedefaultroute

This option is a flag to the *defaultroute* option. If *defaultroute* is set and this flag is also set, *pppd* replaces an existing default route with the new default route. This option is privileged.

disconnect *script*

Execute the command specified by *script*, by passing it to a shell, after *pppd* has terminated the link. This command could, for example, issue commands to the modem to cause it to hang up if hardware modem control signals were not available. The disconnect script is not run if the modem has already hung up. A value for this option from a privileged source cannot be overridden by a non-privileged user.

escape *xx,yy,...*

Specifies that certain characters should be escaped on transmission (regardless of whether the peer requests them to be escaped with its async control character map). The characters to be escaped are specified as a list of hex numbers separated by commas. Note that almost any character can be specified for the *escape* option, unlike the *asynctest* option which only allows control characters to be specified. The characters which may not be escaped are those with hex values 0x20 - 0x3f or 0x5e.

file *name*

Read options from file *name* (the format is described below). The file must be readable by the user who has invoked *pppd*.

init *script*

Execute the command specified by *script*, by passing it to a shell, to initialize the serial line. This script would typically use the *chat(8)* program to configure the modem to enable auto answer. A value for this option from a privileged source cannot be overridden by a non-privileged user.

lock

Specifies that *pppd* should create a UUCP-style lock file for the serial device to ensure exclusive access to the device. By default, *pppd* will not create a lock file.

mru *n*

Set the MRU [Maximum Receive Unit] value to *n*. *Pppd* will ask the peer to send packets of no more than *n* bytes. The value of *n* must be between 128 and 16384; the default is 1500. A value of 296 works well on very slow links (40 bytes for TCP/IP header + 256 bytes of data). Note that for the IPv6 protocol, the MRU must be at least 1280.

mtu *n*

Set the MTU [Maximum Transmit Unit] value to *n*. Unless the peer requests a smaller value via MRU negotiation, *pppd* will request that the kernel networking code send data packets of no more than *n* bytes through the PPP network interface. Note that for the IPv6 protocol, the MTU must be at least 1280.

passive

Enables the "passive" option in the LCP. With this option, *pppd* will attempt to initiate a connection; if no reply is received from the peer, *pppd* will then just wait passively for a valid LCP packet from the peer, instead of exiting, as it would without this option.

OPTIONS

<local_IP_address>:<remote_IP_address>

Set the local and/or remote interface IP addresses. Either one may be omitted. The IP addresses can be specified with a host name or in decimal dot notation (e.g. 150.234.56.78). The default local address is the (first) IP address of the system (unless the *noipdefault* option is given). The remote address will be obtained from the peer if not specified in any option. Thus, in simple cases, this option is not required. If a local and/or remote IP address is specified with this option, *pppd* will not accept a different value from the peer in the IPCP negotiation, unless the *ipcp-accept-local* and/or *ipcp-accept-remote* options are given, respectively.

+ipv6 Enable the IPv6CP and IPv6 protocols.

ipv6 **<local_interface_identifier>,<remote_interface_identifier>**

Set the local and/or remote 64-bit interface identifier. Either one may be omitted. The identifier must be specified in standard ASCII notation of IPv6 addresses (e.g. ::dead:beef). If the *ipv6cp-use-ipaddr* option is given, the local identifier is the local IPv4 address (see above). On systems which supports a unique persistent id, such as EUI-48 derived from the Ethernet MAC address, *ipv6cp-use-persistent* option can be used to replace the *ipv6* **<local>,<remote>** option. Otherwise the identifier is randomized.

active-filter *filter-expression*

Specifies a packet filter to be applied to data packets to determine which packets are to be regarded as link activity, and therefore reset the idle timer, or cause the link to be brought up in demand-dialing mode. This option is useful in conjunction with the **idle** option if there are packets being sent or received regularly over the link (for example, routing information packets) which would otherwise prevent the link from ever appearing to be idle. The *filter-expression* syntax is as described for *tcpdump(1)*, except that qualifiers which are inappropriate for a PPP link, such as **ether** and **arp**, are not permitted. Generally the filter expression should be enclosed in single-quotes to prevent whitespace in the expression from being interpreted by the shell. This option is currently only available under Linux, and requires that the kernel was configured to include PPP filtering support (CONFIG_PPP_FILTER). Note that it is possible to apply different constraints to incoming and outgoing packets using the **inbound** and **outbound** qualifiers.

allow-ip *address(es)*

Allow peers to use the given IP address or subnet without authenticating themselves. The parameter is parsed as for each element of the list of allowed IP addresses in the secrets files (see the AUTHENTICATION section below).

allow-number *number*

Allow peers to connect from the given telephone number. A trailing '*' character will match all numbers beginning with the leading part.

bsdcomp *nr,nt*

Request that the peer compress packets that it sends, using the BSD-Compress scheme, with a maximum code size of *nr* bits, and agree to compress packets sent to the peer with a maximum code size of *nt* bits. If *nt* is not specified, it defaults to the value given for *nr*. Values in the range 9 to 15 may be used for *nr* and *nt*; larger values give better compression but consume more kernel memory for compression dictionaries. Alternatively, a value of 0 for *nr* or *nt* disables compression in the corresponding direction. Use *nobsdcomp* or *bsdcomp 0* to disable BSD-Compress compression entirely.

ca *ca-file*

(EAP-TLS) Use the file *ca-file* as the X.509 Certificate Authority (CA) file (in PEM format), needed for setting up an EAP-TLS connection. This option is used on the client-side in conjunction with the **cert** and **key** options.

cdtrcts Use a non-standard hardware flow control (i.e. DTR/CTS) to control the flow of data on the serial port. If neither *thecrtscts*, the *nocrtscts*, the *cdtr cts* nor the *nocdtrcts* option is given, the hardware flow control setting for the serial port is left unchanged. Some serial ports (such as Macintosh

serial ports) lack a true RTS output. Such serial ports use this mode to implement true bi-directional flow control. The sacrifice is that this flow control mode does not permit using DTR as a modem control line.

cert *certfile*

(EAP-TLS) Use the file *certfile* as the X.509 certificate (in PEM format), needed for setting up an EAP-TLS connection. This option is used on the client-side in conjunction with the **ca** and **key** options.

chap-interval *n*

If this option is given, **pppd** will rechallenge the peer every *n* seconds.

chap-max-challenge *n*

Set the maximum number of CHAP challenge transmissions to *n* (default 10).

chap-restart *n*

Set the CHAP restart interval (retransmission timeout for challenges) to *n* seconds (default 3).

chap-timeout *n*

Set timeout for CHAP authentication by peer to *n* seconds (default 60).

chapms-strip-domain

Some Windows 9x/ME clients might be transmitting the MS domain before the username in the provided client name. This option enables stripping the domain from the client name on the server side before matching it against the secret file.

child-timeout *n*

When exiting, wait for up to *n* seconds for any child processes (such as the command specified with the **pty** command) to exit before exiting. At the end of the timeout, **pppd** will send a SIGTERM signal to any remaining child processes and exit. A value of 0 means no timeout, that is, **pppd** will wait until all child processes have exited.

connect-delay *n*

Wait for up to *n* milliseconds after the connect script finishes for a valid PPP packet from the peer. At the end of this time, or when a valid PPP packet is received from the peer, **pppd** will commence negotiation by sending its first LCP packet. The default value is 1000 (1 second). This wait period only applies if the **connect** or **pty** option is used.

crl *filename*

(EAP-TLS) Use the file *filename* as the Certificate Revocation List to check for the validity of the peer's certificate. This option is not mandatory for setting up an EAP-TLS connection. Also see the **crl-dir** option.

crl-dir *directory*

(EAP-TLS) Use the directory *directory* to scan for CRL files in has format (\$hash.r0) to check for the validity of the peer's certificate. This option is not mandatory for setting up an EAP-TLS connection. Also see the **crl** option.

debug Enables connection debugging facilities. If this option is given, **pppd** will log the contents of all control packets sent or received in a readable form. The packets are logged through syslog with facility *daemon* and level *debug*. This information can be directed to a file by setting up */etc/syslog.conf* appropriately (see *syslog.conf*(5)).

default-asynctest

Disable asynctest negotiation, forcing all control characters to be escaped for both the transmit and the receive direction.

default-mru

Disable MRU [Maximum Receive Unit] negotiation. With this option, **pppd** will use the default MRU value of 1500 bytes for both the transmit and receive direction.

deflate *nr,nt*

Request that the peer compress packets that it sends, using the Deflate scheme, with a maximum window size of 2^{nr} bytes, and agree to compress packets sent to the peer with a maximum window size of 2^{nt} bytes. If *nt* is not specified, it defaults to the value given for *nr*. Values in the range 9 to 15 may be used for *nr* and *nt*; larger values give better compression but consume more kernel memory for compression dictionaries. Alternatively, a value of 0 for *nr* or *nt* disables compression in the corresponding direction. Use *nodeflate* or *deflate 0* to disable Deflate compression entirely. (Note: *pppd* requests Deflate compression in preference to BSD-Compress if the peer can do either.)

demand

Initiate the link only on demand, i.e. when data traffic is present. With this option, the remote IP address may be specified by the user on the command line or in an options file, or if not, *pppd* will use an arbitrary address in the 10.x.x.x range. *Pppd* will initially configure the interface and enable it for IP traffic without connecting to the peer. When traffic is available, *pppd* will connect to the peer and perform negotiation, authentication, etc. When this is completed, *pppd* will commence passing data packets (i.e., IP packets) across the link.

The *demand* option implies the *persist* option. If this behaviour is not desired, use the *nopersist* option after the *demand* option. The *idle* and *holdoff* options are also useful in conjunction with the *demand* option.

domain *d*

Append the domain name *d* to the local host name for authentication purposes. For example, if *gethostname()* returns the name *porsche*, but the fully qualified domain name is *porsche.Quotron.COM*, you could specify *domain Quotron.COM*. *Pppd* would then use the name *porsche.Quotron.COM* for looking up secrets in the secrets file, and as the default name to send to the peer when authenticating itself to the peer. This option is privileged.

dryrun With the **dryrun** option, *pppd* will print out all the option values which have been set and then exit, after parsing the command line and options files and checking the option values, but before initiating the link. The option values are logged at level *info*, and also printed to standard output unless the device on standard output is the device that *pppd* would be using to communicate with the peer.

dump With the **dump** option, *pppd* will print out all the option values which have been set. This option is like the **dryrun** option except that *pppd* proceeds as normal rather than exiting.

enable-session

Enables session accounting via PAM or *wtmp/wtmpx*, as appropriate. When PAM is enabled, the PAM "account" and "session" module stacks determine behavior, and are enabled for all PPP authentication protocols. When PAM is disabled, *wtmp/wtmpx* entries are recorded regardless of whether the peer name identifies a valid user on the local system, making peers visible in the *last(1)* log. This feature is automatically enabled when the *pppd login* option is used. Session accounting is disabled by default.

endpoint *<epdisc>*

Sets the endpoint discriminator sent by the local machine to the peer during multilink negotiation to *<epdisc>*. The default is to use the MAC address of the first ethernet interface on the system, if any, otherwise the IPv4 address corresponding to the hostname, if any, provided it is not in the multicast or locally-assigned IP address ranges, or the localhost address. The endpoint discriminator can be the string **null** or of the form *type:value*, where *type* is a decimal number or one of the strings **local**, **IP**, **MAC**, **magic**, or **phone**. The value is an IP address in dotted-decimal notation for the **IP** type, or a string of bytes in hexadecimal, separated by periods or colons for the other types. For the **MAC** type, the value may also be the name of an ethernet or similar network interface. This option is currently only available under Linux.

eap-interval *n*

If this option is given and `pppd` authenticates the peer with EAP (i.e., is the server), `pppd` will restart EAP authentication every *n* seconds. For EAP SRP-SHA1, see also the **srp-interval** option, which enables lightweight rechallenge.

eap-max-rreq *n*

Set the maximum number of EAP Requests to which `pppd` will respond (as a client) without hearing EAP Success or Failure. (Default is 20.)

eap-max-sreq *n*

Set the maximum number of EAP Requests that `pppd` will issue (as a server) while attempting authentication. (Default is 10.)

eap-restart *n*

Set the retransmit timeout for EAP Requests when acting as a server (authenticator). (Default is 3 seconds.)

eap-timeout *n*

Set the maximum time to wait for the peer to send an EAP Request when acting as a client (authenticatee). (Default is 20 seconds.)

hide-password

When logging the contents of PAP packets, this option causes `pppd` to exclude the password string from the log. This is the default.

holdoff *n*

Specifies how many seconds to wait before re-initiating the link after it terminates. This option only has any effect if the *persist* or *demand* option is used. The holdoff period is not applied if the link was terminated because it was idle.

idle *n* Specifies that `pppd` should disconnect if the link is idle for *n* seconds. The link is idle when no data packets (i.e. IP packets) are being sent or received. Note: it is not advisable to use this option with the *persist* option without the *demand* option. If the **active-filter** option is given, data packets which are rejected by the specified activity filter also count as the link being idle.

ipcp-accept-local

With this option, `pppd` will accept the peer's idea of our local IP address, even if the local IP address was specified in an option.

ipcp-accept-remote

With this option, `pppd` will accept the peer's idea of its (remote) IP address, even if the remote IP address was specified in an option.

ipcp-max-configure *n*

Set the maximum number of IPCP configure-request transmissions to *n* (default 10).

ipcp-max-failure *n*

Set the maximum number of IPCP configure-NAKs returned before starting to send configure-Rejects instead to *n* (default 10).

ipcp-max-terminate *n*

Set the maximum number of IPCP terminate-request transmissions to *n* (default 3).

ipcp-restart *n*

Set the IPCP restart interval (retransmission timeout) to *n* seconds (default 3).

ipparam *string*

Provides an extra parameter to the `ip-up`, `ip-pre-up` and `ip-down` scripts. If this option is given, the *string* supplied is given as the 6th parameter to those scripts.

ipv6cp-accept-local

With this option, `pppd` will accept the peer's idea of our local IPv6 interface identifier, even if the local IPv6 interface identifier was specified in an option.

ipv6cp-accept-remote

With this option, pppd will accept the peer's idea of its (remote) IPv6 interface identifier, even if the remote IPv6 interface identifier was specified in an option.

ipv6cp-max-configure *n*

Set the maximum number of IPv6CP configure-request transmissions to *n* (default 10).

ipv6cp-max-failure *n*

Set the maximum number of IPv6CP configure-NAKs returned before starting to send configure-Rejects instead to *n* (default 10).

ipv6cp-max-terminate *n*

Set the maximum number of IPv6CP terminate-request transmissions to *n* (default 3).

ipv6cp-restart *n*

Set the IPv6CP restart interval (retransmission timeout) to *n* seconds (default 3).

ipx

Enable the IPXCP and IPX protocols. This option is presently only supported under Linux, and only if your kernel has been configured to include IPX support.

ipx-network *n*

Set the IPX network number in the IPXCP configure request frame to *n*, a hexadecimal number (without a leading 0x). There is no valid default. If this option is not specified, the network number is obtained from the peer. If the peer does not have the network number, the IPX protocol will not be started.

ipx-node *n:m*

Set the IPX node numbers. The two node numbers are separated from each other with a colon character. The first number *n* is the local node number. The second number *m* is the peer's node number. Each node number is a hexadecimal number, at most 10 digits long. The node numbers on the ipx-network must be unique. There is no valid default. If this option is not specified then the node numbers are obtained from the peer.

ipx-router-name *<string>*

Set the name of the router. This is a string and is sent to the peer as information data.

ipx-routing *n*

Set the routing protocol to be received by this option. More than one instance of *ipx-routing* may be specified. The 'none' option (0) may be specified as the only instance of ipx-routing. The values may be 0 for *NONE*, 2 for *RIP/SAP*, and 4 for *NLSP*.

ipxcp-accept-local

Accept the peer's NAK for the node number specified in the ipx-node option. If a node number was specified, and non-zero, the default is to insist that the value be used. If you include this option then you will permit the peer to override the entry of the node number.

ipxcp-accept-network

Accept the peer's NAK for the network number specified in the ipx-network option. If a network number was specified, and non-zero, the default is to insist that the value be used. If you include this option then you will permit the peer to override the entry of the node number.

ipxcp-accept-remote

Use the peer's network number specified in the configure request frame. If a node number was specified for the peer and this option was not specified, the peer will be forced to use the value which you have specified.

ipxcp-max-configure *n*

Set the maximum number of IPXCP configure request frames which the system will send to *n*. The default is 10.

ipxcp-max-failure *n*

Set the maximum number of IPXCP NAK frames which the local system will send before it rejects the options. The default value is 3.

ipxcp-max-terminate *n*

Set the maximum number of IPXCP terminate request frames before the local system considers that the peer is not listening to them. The default value is 3.

kdebug *n*

Enable debugging code in the kernel-level PPP driver. The argument values depend on the specific kernel driver, but in general a value of 1 will enable general kernel debug messages. (Note that these messages are usually only useful for debugging the kernel driver itself.) For the Linux 2.2.x kernel driver, the value is a sum of bits: 1 to enable general debug messages, 2 to request that the contents of received packets be printed, and 4 to request that the contents of transmitted packets be printed. On most systems, messages printed by the kernel are logged by syslog(1) to a file as directed in the `/etc/syslog.conf` configuration file.

key *keyfile*

(EAP-TLS) Use the file *keyfile* as the private key file (in PEM format), needed for setting up an EAP-TLS connection. This option is used on the client-side in conjunction with the **ca** and **cert** options.

ktune Enables pppd to alter kernel settings as appropriate. Under Linux, pppd will enable IP forwarding (i.e. set `/proc/sys/net/ipv4/ip_forward` to 1) if the *proxyarp* option is used, and will enable the dynamic IP address option (i.e. set `/proc/sys/net/ipv4/ip_dynaddr` to 1) in demand mode if the local address changes.

lcp-echo-adaptive

If this option is used with the *lcp-echo-failure* option then pppd will send LCP echo-request frames only if no traffic was received from the peer since the last echo-request was sent.

lcp-echo-failure *n*

If this option is given, pppd will presume the peer to be dead if *n* LCP echo-requests are sent without receiving a valid LCP echo-reply. If this happens, pppd will terminate the connection. Use of this option requires a non-zero value for the *lcp-echo-interval* parameter. This option can be used to enable pppd to terminate after the physical connection has been broken (e.g., the modem has hung up) in situations where no hardware modem control lines are available.

lcp-echo-interval *n*

If this option is given, pppd will send an LCP echo-request frame to the peer every *n* seconds. Normally the peer should respond to the echo-request by sending an echo-reply. This option can be used with the *lcp-echo-failure* option to detect that the peer is no longer connected.

lcp-max-configure *n*

Set the maximum number of LCP configure-request transmissions to *n* (default 10).

lcp-max-failure *n*

Set the maximum number of LCP configure-NAKs returned before starting to send configure-Rejects instead to *n* (default 10).

lcp-max-terminate *n*

Set the maximum number of LCP terminate-request transmissions to *n* (default 3).

lcp-restart *n*

Set the LCP restart interval (retransmission timeout) to *n* seconds (default 3).

linkname *name*

Sets the logical name of the link to *name*. Pppd will create a file named **ppp-name.pid** in `/var/run` (or `/etc/ppp` on some systems) containing its process ID. This can be useful in determining which instance of pppd is responsible for the link to a given peer system. This is a privileged option.

local Don't use the modem control lines. With this option, pppd will ignore the state of the CD (Carrier Detect) signal from the modem and will not change the state of the DTR (Data Terminal Ready) signal. This is the opposite of the **modem** option.

logfd *n* Send log messages to file descriptor *n*. Pppd will send log messages to at most one file or file descriptor (as well as sending the log messages to syslog), so this option and the **logfile** option are mutually exclusive. The default is for pppd to send log messages to stdout (file descriptor 1), unless the serial port is already open on stdout.

logfile *filename*

Append log messages to the file *filename* (as well as sending the log messages to syslog). The file is opened with the privileges of the user who invoked pppd, in append mode.

login Use the system password database for authenticating the peer using PAP, and record the user in the system wtmp file. Note that the peer must have an entry in the /etc/ppp/pap-secrets file as well as the system password database to be allowed access. See also the **enable-session** option.

master_detach

If multilink is enabled and this pppd process is the multilink bundle master, and the link controlled by this pppd process terminates, this pppd process continues to run in order to maintain the bundle. If the **master_detach** option has been given, pppd will detach from its controlling terminal in this situation, even if the **nodetach** option has been given.

maxconnect *n*

Terminate the connection when it has been available for network traffic for *n* seconds (i.e. *n* seconds after the first network control protocol comes up).

maxfail *n*

Terminate after *n* consecutive failed connection attempts. A value of 0 means no limit. The default value is 10.

modem

Use the modem control lines. This option is the default. With this option, pppd will wait for the CD (Carrier Detect) signal from the modem to be asserted when opening the serial device (unless a connect script is specified), and it will drop the DTR (Data Terminal Ready) signal briefly when the connection is terminated and before executing the connect script. On Ultrix, this option implies hardware flow control, as for the **crtsets** option. This is the opposite of the **local** option.

mp Enables the use of PPP multilink; this is an alias for the 'multilink' option. This option is currently only available under Linux.

mppe-stateful

Allow MPPE to use stateful mode. Stateless mode is still attempted first. The default is to disallow stateful mode.

mpshortseq

Enables the use of short (12-bit) sequence numbers in multilink headers, as opposed to 24-bit sequence numbers. This option is only available under Linux, and only has any effect if multilink is enabled (see the multilink option).

mrru *n* Sets the Maximum Reconstructed Receive Unit to *n*. The MRRU is the maximum size for a received packet on a multilink bundle, and is analogous to the MRU for the individual links. This option is currently only available under Linux, and only has any effect if multilink is enabled (see the multilink option).

ms-dns <*addr*>

If pppd is acting as a server for Microsoft Windows clients, this option allows pppd to supply one or two DNS (Domain Name Server) addresses to the clients. The first instance of this option specifies the primary DNS address; the second instance (if given) specifies the secondary DNS address. (This option was present in some older versions of pppd under the name **dns-addr**.)

ms-wins <*addr*>

If pppd is acting as a server for Microsoft Windows or "Samba" clients, this option allows pppd to supply one or two WINS (Windows Internet Name Services) server addresses to the clients. The first instance of this option specifies the primary WINS address; the second instance (if given) specifies the secondary WINS address.

multilink

Enables the use of the PPP multilink protocol. If the peer also supports multilink, then this link can become part of a bundle between the local system and the peer. If there is an existing bundle to the peer, `pppd` will join this link to that bundle, otherwise `pppd` will create a new bundle. See the MULTILINK section below. This option is currently only available under Linux.

name *name*

Set the name of the local system for authentication purposes to *name*. This is a privileged option. With this option, `pppd` will use lines in the secrets files which have *name* as the second field when looking for a secret to use in authenticating the peer. In addition, unless overridden with the *user* option, *name* will be used as the name to send to the peer when authenticating the local system to the peer. (Note that `pppd` does not append the domain name to *name*.)

noaccomp

Disable Address/Control compression in both directions (send and receive).

need-peer-eap

(EAP-TLS) Require the peer to verify our authentication credentials.

noauth Do not require the peer to authenticate itself. This option is privileged.

nobsdcomp

Disables BSD-Compress compression; **pppd** will not request or agree to compress packets using the BSD-Compress scheme.

noccp Disable CCP (Compression Control Protocol) negotiation. This option should only be required if the peer is buggy and gets confused by requests from `pppd` for CCP negotiation.

nocrtscts

Disable hardware flow control (i.e. RTS/CTS) on the serial port. If neither the *crtcts* nor the *nocrtcts* nor the *cdtrcts* nor the *nocdtrcts* option is given, the hardware flow control setting for the serial port is left unchanged.

nocdtrcts

This option is a synonym for *nocrtcts*. Either of these options will disable both forms of hardware flow control.

nodefaultroute

Disable the *defaultroute* option. The system administrator who wishes to prevent users from adding a default route with `pppd` can do so by placing this option in the `/etc/ppp/options` file.

noreplacedefaultroute

Disable the *replacedefaultroute* option. This allows to disable a *replacedefaultroute* option set previously in the configuration.

nodefaultroute6

Disable the *defaultroute6* option. The system administrator who wishes to prevent users from adding a default route with `pppd` can do so by placing this option in the `/etc/ppp/options` file.

nodeflate

Disables Deflate compression; `pppd` will not request or agree to compress packets using the Deflate scheme.

nodetach

Don't detach from the controlling terminal. Without this option, if a serial device other than the terminal on the standard input is specified, `pppd` will fork to become a background process.

noendpoint

Disables `pppd` from sending an endpoint discriminator to the peer or accepting one from the peer (see the MULTILINK section below). This option should only be required if the peer is buggy.

noip

Disable IPCP negotiation and IP communication. This option should only be required if the peer is buggy and gets confused by requests from `pppd` for IPCP negotiation.

- noipv6** Disable IPv6CP negotiation and IPv6 communication. This option should only be required if the peer is buggy and gets confused by requests from pppd for IPv6CP negotiation.
- noipdefault**
Disables the default behaviour when no local IP address is specified, which is to determine (if possible) the local IP address from the hostname. With this option, the peer will have to supply the local IP address during IPCP negotiation (unless it specified explicitly on the command line or in an options file).
- noipx** Disable the IPXCP and IPX protocols. This option should only be required if the peer is buggy and gets confused by requests from pppd for IPXCP negotiation.
- noktune**
Opposite of the *ktune* option; disables pppd from changing system settings.
- nolock** Opposite of the *lock* option; specifies that pppd should not create a UUCP-style lock file for the serial device. This option is privileged.
- nolog** Do not send log messages to a file or file descriptor. This option cancels the **logfd** and **logfile** options.
- nomagic**
Disable magic number negotiation. With this option, pppd cannot detect a looped-back line. This option should only be needed if the peer is buggy.
- nomp** Disables the use of PPP multilink. This option is currently only available under Linux.
- nomppe**
Disables MPPE (Microsoft Point to Point Encryption). This is the default.
- nomppe-40**
Disable 40-bit encryption with MPPE.
- nomppe-128**
Disable 128-bit encryption with MPPE.
- nomppe-stateful**
Disable MPPE stateful mode. This is the default.
- nompshortseq**
Disables the use of short (12-bit) sequence numbers in the PPP multilink protocol, forcing the use of 24-bit sequence numbers. This option is currently only available under Linux, and only has any effect if multilink is enabled.
- nomultilink**
Disables the use of PPP multilink. This option is currently only available under Linux.
- nopcomp**
Disable protocol field compression negotiation in both the receive and the transmit direction.
- nopersist**
Exit once a connection has been made and terminated. This is the default unless the *persist* or *demand* option has been specified.
- nopredictor1**
Do not accept or agree to Predictor-1 compression.
- noproxyarp**
Disable the *proxyarp* option. The system administrator who wishes to prevent users from creating proxy ARP entries with pppd can do so by placing this option in the */etc/ppp/options* file.
- noremoveip**
Allow pppd to operate without having an IP address for the peer. This option is only available under Linux. Normally, pppd will request the peer's IP address, and if the peer does not supply it, pppd will use an arbitrary address in the 10.x.x.x subnet. With this option, if the peer does not

supply its IP address, `pppd` will not ask the peer for it, and will not set the destination address of the ppp interface. In this situation, the ppp interface can be used for routing by creating device routes, but the peer itself cannot be addressed directly for IP traffic.

notty Normally, `pppd` requires a terminal device. With this option, `pppd` will allocate itself a pseudo-tty master/slave pair and use the slave as its terminal device. `pppd` will create a child process to act as a 'character shunt' to transfer characters between the pseudo-tty master and its standard input and output. Thus `pppd` will transmit characters on its standard output and receive characters on its standard input even if they are not terminal devices. This option increases the latency and CPU overhead of transferring data over the ppp interface as all of the characters sent and received must flow through the character shunt process. An explicit device name may not be given if this option is used.

novj Disable Van Jacobson style TCP/IP header compression in both the transmit and the receive direction.

novjccomp

Disable the connection-ID compression option in Van Jacobson style TCP/IP header compression. With this option, `pppd` will not omit the connection-ID byte from Van Jacobson compressed TCP/IP headers, nor ask the peer to do so.

papcrypt

Indicates that all secrets in the `/etc/ppp/pap-secrets` file which are used for checking the identity of the peer are encrypted, and thus `pppd` should not accept a password which, before encryption, is identical to the secret from the `/etc/ppp/pap-secrets` file.

pap-max-authreq *n*

Set the maximum number of PAP authenticate-request transmissions to *n* (default 10).

pap-restart *n*

Set the PAP restart interval (retransmission timeout) to *n* seconds (default 3).

pap-timeout *n*

Set the maximum time that `pppd` will wait for the peer to authenticate itself with PAP to *n* seconds (0 means no limit).

pass-filter *filter-expression*

Specifies a packet filter to be applied to data packets being sent or received to determine which packets should be allowed to pass. Packets which are rejected by the filter are silently discarded. This option can be used to prevent specific network daemons (such as routed) using up link bandwidth, or to provide a very basic firewall capability. The *filter-expression* syntax is as described for `tcpdump(1)`, except that qualifiers which are inappropriate for a PPP link, such as **ether** and **arp**, are not permitted. Generally the filter expression should be enclosed in single-quotes to prevent whitespace in the expression from being interpreted by the shell. Note that it is possible to apply different constraints to incoming and outgoing packets using the **inbound** and **outbound** qualifiers. This option is currently only available under Linux, and requires that the kernel was configured to include PPP filtering support (`CONFIG_PPP_FILTER`).

password *password-string*

Specifies the password to use for authenticating to the peer. Use of this option is discouraged, as the password is likely to be visible to other users on the system (for example, by using `ps(1)`).

persist Do not exit after a connection is terminated; instead try to reopen the connection. The **maxfail** option still has an effect on persistent connections.

plugin *filename*

Load the shared library object file *filename* as a plugin. This is a privileged option. If *filename* does not contain a slash (/), `pppd` will look in the `/usr/lib/pppd/version` directory for the plugin, where *version* is the version number of `pppd` (for example, 2.4.2).

predictor1

Request that the peer compress frames that it sends using Predictor-1 compression, and agree to compress transmitted frames with Predictor-1 if requested. This option has no effect unless the kernel driver supports Predictor-1 compression.

privgroup *group-name*

Allows members of group *group-name* to use privileged options. This is a privileged option. Use of this option requires care as there is no guarantee that members of *group-name* cannot use pppd to become root themselves. Consider it equivalent to putting the members of *group-name* in the kmem or disk group.

proxyarp

Add an entry to this system's ARP [Address Resolution Protocol] table with the IP address of the peer and the Ethernet address of this system. This will have the effect of making the peer appear to other systems to be on the local ethernet.

pty *script*

Specifies that the command *script* is to be used to communicate rather than a specific terminal device. Pppd will allocate itself a pseudo-tty master/slave pair and use the slave as its terminal device. The *script* will be run in a child process with the pseudo-tty master as its standard input and output. An explicit device name may not be given if this option is used. (Note: if the *record* option is used in conjunction with the *pty* option, the child process will have pipes on its standard input and output.)

receive-all

With this option, pppd will accept all control characters from the peer, including those marked in the receive asynmap. Without this option, pppd will discard those characters as specified in RFC1662. This option should only be needed if the peer is buggy.

record *filename*

Specifies that pppd should record all characters sent and received to a file named *filename*. This file is opened in append mode, using the user's user-ID and permissions. This option is implemented using a pseudo-tty and a process to transfer characters between the pseudo-tty and the real serial device, so it will increase the latency and CPU overhead of transferring data over the ppp interface. The characters are stored in a tagged format with timestamps, which can be displayed in readable form using the pppdump(8) program.

remotename *name*

Set the assumed name of the remote system for authentication purposes to *name*.

remotenum *number*

Set the assumed telephone number of the remote system for authentication purposes to *number*.

refuse-chap

With this option, pppd will not agree to authenticate itself to the peer using CHAP.

refuse-mschap

With this option, pppd will not agree to authenticate itself to the peer using MS-CHAP.

refuse-mschap-v2

With this option, pppd will not agree to authenticate itself to the peer using MS-CHAPv2.

refuse-eap

With this option, pppd will not agree to authenticate itself to the peer using EAP.

refuse-pap

With this option, pppd will not agree to authenticate itself to the peer using PAP.

require-chap

Require the peer to authenticate itself using CHAP [Challenge Handshake Authentication Protocol] authentication.

require-mppe

Require the use of MPPE (Microsoft Point to Point Encryption). This option disables all other compression types. This option enables both 40-bit and 128-bit encryption. In order for MPPE to successfully come up, you must have authenticated with either MS-CHAP or MS-CHAPv2. This option is presently only supported under Linux, and only if your kernel has been configured to include MPPE support.

require-mppe-40

Require the use of MPPE, with 40-bit encryption.

require-mppe-128

Require the use of MPPE, with 128-bit encryption.

require-mschap

Require the peer to authenticate itself using MS-CHAP [Microsoft Challenge Handshake Authentication Protocol] authentication.

require-mschap-v2

Require the peer to authenticate itself using MS-CHAPv2 [Microsoft Challenge Handshake Authentication Protocol, Version 2] authentication.

require-eap

Require the peer to authenticate itself using EAP [Extensible Authentication Protocol] authentication.

require-pap

Require the peer to authenticate itself using PAP [Password Authentication Protocol] authentication.

set name=value

Set an environment variable for scripts that are invoked by `pppd`. When set by a privileged source, the variable specified by *name* cannot be changed by options contained in an unprivileged source. See also the *unset* option and the environment described in *SCRIPTS*.

show-password

When logging the contents of PAP packets, this option causes `pppd` to show the password string in the log message.

silent With this option, `pppd` will not transmit LCP packets to initiate a connection until a valid LCP packet is received from the peer (as for the 'passive' option with ancient versions of `pppd`).

srp-interval *n*

If this parameter is given and `pppd` uses EAP SRP-SHA1 to authenticate the peer (i.e., is the server), then `pppd` will use the optional lightweight SRP rechallenge mechanism at intervals of *n* seconds. This option is faster than **eap-interval** reauthentication because it uses a hash-based mechanism and does not derive a new session key.

srp-pn-secret *string*

Set the long-term pseudonym-generating secret for the server. This value is optional and if set, needs to be known at the server (authenticator) side only, and should be different for each server (or pool of identical servers). It is used along with the current date to generate a key to encrypt and decrypt the client's identity contained in the pseudonym.

srp-use-pseudonym

When operating as an EAP SRP-SHA1 client, attempt to use the pseudonym stored in `~/ppp_pseudonym` first as the identity, and save in this file any pseudonym offered by the peer during authentication.

stop-bits *n*

Set the number of stop bits for the serial port. Valid values are 1 or 2. The default value is 1.

sync

Use synchronous HDLC serial encoding instead of asynchronous. The device used by `pppd` with this option must have sync support. Currently supports Microgate SyncLink adapters under Linux

and FreeBSD 2.2.8 and later.

unit *num*

Sets the ppp unit number (for a ppp0 or ppp1 etc interface name) for outbound connections. If the unit is already in use a dynamically allocated number will be used.

ifname *string*

Set the ppp interface name for outbound connections. If the interface name is already in use, or if the name cannot be used for any other reason, pppd will terminate.

unset *name*

Remove a variable from the environment variable for scripts that are invoked by pppd. When specified by a privileged source, the variable *name* cannot be set by options contained in an unprivileged source. See also the *set* option and the environment described in *SCRIPTS*.

updetach

With this option, pppd will detach from its controlling terminal once it has successfully established the ppp connection (to the point where the first network control protocol, usually the IP control protocol, has come up).

up_sdnotify

Use this option to run pppd in systemd service units of Type=notify (**up_sdnotify** implies **node-tach**). When **up_sdnotify** is enabled, pppd will notify systemd once it has successfully established the ppp connection (to the point where the first network control protocol, usually the IP control protocol, has come up). This option is only available when pppd is compiled with systemd support.

usehostname

Enforce the use of the hostname (with domain name appended, if given) as the name of the local system for authentication purposes (overrides the *name* option). This option is not normally needed since the *name* option is privileged.

usepeerdns

Ask the peer for up to 2 DNS server addresses. The addresses supplied by the peer (if any) are passed to the /etc/ppp/ip-up script in the environment variables DNS1 and DNS2, and the environment variable USEPEERDNS will be set to 1. In addition, pppd will create an /etc/ppp/resolv.conf file containing one or two nameserver lines with the address(es) supplied by the peer.

user *name*

Sets the name used for authenticating the local system to the peer to *name*.

vj-max-slots *n*

Sets the number of connection slots to be used by the Van Jacobson TCP/IP header compression and decompression code to *n*, which must be between 2 and 16 (inclusive).

welcome *script*

Run the executable or shell command specified by *script* before initiating PPP negotiation, after the connect script (if any) has completed. A value for this option from a privileged source cannot be overridden by a non-privileged user.

xonxoff

Use software flow control (i.e. XON/XOFF) to control the flow of data on the serial port.

PPPOE OPTIONS

To establish PPP link over Ethernet (PPPoE) it is needed to load pppd's **plugin pppoe.so** and then specify option **nic-interface** instead of modem options *ttyname* and *speed*. Recognized pppd's PPPoE options are:

nic-interface

Use the ethernet device *interface* to communicate with the peer. For example, establishing PPPoE link on *eth0* interface is done by specifying pppd option **nic-eth0**. Prefix **nic-** for this option may be avoided if interface name is unambiguous and does not look like any other pppd's option.

pppoe-service *name*

Connect to specified PPPoE service name. For backward compatibility also **rp_pppoe_service** option name is supported.

pppoe-ac *name*

Connect to specified PPPoE access concentrator name. For backward compatibility also **rp_pppoe_ac** option name is supported.

pppoe-sess *sessid:macaddr*

Attach to existing PPPoE session. For backward compatibility also **rp_pppoe_sess** option name is supported.

pppoe-verbose *n*

Be verbose about discovered access concentrators. For backward compatibility also **rp_pppoe_verbose** option name is supported.

pppoe-mac *macaddr*

Connect to specified MAC address.

pppoe-host-uniq *string*

Set the PPPoE Host-Uniq tag to the supplied hex string. By default PPPoE Host-Uniq tag is set to the pppd's process PID. For backward compatibility this option may be specified without **pppoe-**prefix.

pppoe-padi-timeout *n*

Initial timeout for discovery packets in seconds (default 5).

pppoe-padi-attempts *n*

Number of discovery attempts (default 3).

OPTIONS FILES

Options can be taken from files as well as the command line. Pppd reads options from the files */etc/ppp/options*, *~/.ppprc* and */etc/ppp/options.ttyname* (in that order) before processing the options on the command line. (In fact, the command-line options are scanned to find the terminal name before the *options.ttyname* file is read.) In forming the name of the *options.ttyname* file, the initial */dev/* is removed from the terminal name, and any remaining */* characters are replaced with dots.

An options file is parsed into a series of words, delimited by whitespace. Whitespace can be included in a word by enclosing the word in double-quotes (*"*). A backslash (**) quotes the following character. A hash (*#*) starts a comment, which continues until the end of the line. There is no restriction on using the *file* or *call* options within an options file.

SECURITY

pppd provides system administrators with sufficient access control that PPP access to a server machine can be provided to legitimate users without fear of compromising the security of the server or the network it's on. This control is provided through restrictions on which IP addresses the peer may use, based on its authenticated identity (if any), and through restrictions on which options a non-privileged user may use. Several of pppd's options are privileged, in particular those which permit potentially insecure configurations; these options are only accepted in files which are under the control of the system administrator, or if pppd is being run by root.

The default behaviour of pppd is to allow an unauthenticated peer to use a given IP address only if the system does not already have a route to that IP address. For example, a system with a permanent connection to the wider internet will normally have a default route, and thus all peers will have to authenticate themselves in order to set up a connection. On such a system, the *auth* option is the default. On the other hand, a system where the PPP link is the only connection to the internet will not normally have a default route, so the peer will be able to use almost any IP address without authenticating itself.

As indicated above, some security-sensitive options are privileged, which means that they may not be used by an ordinary non-privileged user running a *setuid-root pppd*, either on the command line, in the user's *~/.ppprc* file, or in an options file read using the *file* option. Privileged options may be used in */etc/ppp/options* file or in an options file read using the *call* option. If pppd is being run by the root user, privileged

options can be used without restriction.

When opening the device, `pppd` uses either the invoking user's user ID or the root UID (that is, 0), depending on whether the device name was specified by the user or the system administrator. If the device name comes from a privileged source, that is, `/etc/ppp/options` or an options file read using the `call` option, `pppd` uses full root privileges when opening the device. Thus, by creating an appropriate file under `/etc/ppp/peers`, the system administrator can allow users to establish a ppp connection via a device which they would not normally have permission to access. Otherwise `pppd` uses the invoking user's real UID when opening the device.

AUTHENTICATION

Authentication is the process whereby one peer convinces the other of its identity. This involves the first peer sending its name to the other, together with some kind of secret information which could only come from the genuine authorized user of that name. In such an exchange, we will call the first peer the "client" and the other the "server". The client has a name by which it identifies itself to the server, and the server also has a name by which it identifies itself to the client. Generally the genuine client shares some secret (or password) with the server, and authenticates itself by proving that it knows that secret. Very often, the names used for authentication correspond to the internet hostnames of the peers, but this is not essential.

At present, `pppd` supports three authentication protocols: the Password Authentication Protocol (PAP), Challenge Handshake Authentication Protocol (CHAP), and Extensible Authentication Protocol (EAP). PAP involves the client sending its name and a cleartext password to the server to authenticate itself. In contrast, the server initiates the CHAP authentication exchange by sending a challenge to the client (the challenge packet includes the server's name). The client must respond with a response which includes its name plus a hash value derived from the shared secret and the challenge, in order to prove that it knows the secret. EAP supports CHAP-style authentication, and also includes the SRP-SHA1 mechanism, which is resistant to dictionary-based attacks and does not require a cleartext password on the server side.

The PPP protocol, being symmetrical, allows both peers to require the other to authenticate itself. In that case, two separate and independent authentication exchanges will occur. The two exchanges could use different authentication protocols, and in principle, different names could be used in the two exchanges.

The default behaviour of `pppd` is to agree to authenticate if requested, and to not require authentication from the peer. However, `pppd` will not agree to authenticate itself with a particular protocol if it has no secrets which could be used to do so.

`pppd` stores secrets for use in authentication in secrets files (`/etc/ppp/pap-secrets` for PAP, `/etc/ppp/chap-secrets` for CHAP, MS-CHAP, MS-CHAPv2, and EAP MD5-Challenge, and `/etc/ppp/srp-secrets` for EAP SRP-SHA1). All secrets files have the same format. The secrets files can contain secrets for `pppd` to use in authenticating itself to other systems, as well as secrets for `pppd` to use when authenticating other systems to itself.

Each line in a secrets file contains one secret. A given secret is specific to a particular combination of client and server - it can only be used by that client to authenticate itself to that server. Thus each line in a secrets file has at least 3 fields: the name of the client, the name of the server, and the secret. These fields may be followed by a list of the IP addresses that the specified client may use when connecting to the specified server.

A secrets file is parsed into words as for a options file, so the client name, server name and secrets fields must each be one word, with any embedded spaces or other special characters quoted or escaped. Note that case is significant in the client and server names and in the secret.

If the secret starts with an '@', what follows is assumed to be the name of a file from which to read the secret. A "*" as the client or server name matches any name. When selecting a secret, `pppd` takes the best match, i.e. the match with the fewest wildcards.

Any following words on the same line are taken to be a list of acceptable IP addresses for that client. If there are only 3 words on the line, or if the first word is "-", then all IP addresses are disallowed. To allow any address, use "*". A word starting with "!" indicates that the specified address is *not* acceptable. An address may be followed by "/" and a number *n*, to indicate a whole subnet, i.e. all addresses which have

the same value in the most significant n bits. In this form, the address may be followed by a plus sign ("+") to indicate that one address from the subnet is authorized, based on the ppp network interface unit number in use. In this case, the host part of the address will be set to the unit number plus one.

Thus a secrets file contains both secrets for use in authenticating other hosts, plus secrets which we use for authenticating ourselves to others. When pppd is authenticating the peer (checking the peer's identity), it chooses a secret with the peer's name in the first field and the name of the local system in the second field. The name of the local system defaults to the hostname, with the domain name appended if the *domain* option is used. This default can be overridden with the *name* option, except when the *usehostname* option is used. (For EAP SRP-SHA1, see the *srp-entry(8)* utility for generating proper validator entries to be used in the "secret" field.)

When pppd is choosing a secret to use in authenticating itself to the peer, it first determines what name it is going to use to identify itself to the peer. This name can be specified by the user with the *user* option. If this option is not used, the name defaults to the name of the local system, determined as described in the previous paragraph. Then pppd looks for a secret with this name in the first field and the peer's name in the second field. Pppd will know the name of the peer if CHAP or EAP authentication is being used, because the peer will have sent it in the challenge packet. However, if PAP is being used, pppd will have to determine the peer's name from the options specified by the user. The user can specify the peer's name directly with the *remotename* option. Otherwise, if the remote IP address was specified by a name (rather than in numeric form), that name will be used as the peer's name. Failing that, pppd will use the null string as the peer's name.

When authenticating the peer with PAP, the supplied password is first compared with the secret from the secrets file. If the password doesn't match the secret, the password is encrypted using *crypt()* and checked against the secret again. Thus secrets for authenticating the peer can be stored in encrypted form if desired. If the *papcrypt* option is given, the first (unencrypted) comparison is omitted, for better security.

Furthermore, if the *login* option was specified, the username and password are also checked against the system password database. Thus, the system administrator can set up the *pap-secrets* file to allow PPP access only to certain users, and to restrict the set of IP addresses that each user can use. Typically, when using the *login* option, the secret in */etc/ppp/pap-secrets* would be "", which will match any password supplied by the peer. This avoids the need to have the same secret in two places.

Authentication must be satisfactorily completed before IPCP (or any other Network Control Protocol) can be started. If the peer is required to authenticate itself, and fails to do so, pppd will terminate the link (by closing LCP). If IPCP negotiates an unacceptable IP address for the remote host, IPCP will be closed. IP packets can only be sent or received when IPCP is open.

In some cases it is desirable to allow some hosts which can't authenticate themselves to connect and use one of a restricted set of IP addresses, even when the local host generally requires authentication. If the peer refuses to authenticate itself when requested, pppd takes that as equivalent to authenticating with PAP using the empty string for the username and password. Thus, by adding a line to the *pap-secrets* file which specifies the empty string for the client and password, it is possible to allow restricted access to hosts which refuse to authenticate themselves.

ROUTING

When IPCP negotiation is completed successfully, pppd will inform the kernel of the local and remote IP addresses for the ppp interface. This is sufficient to create a host route to the remote end of the link, which will enable the peers to exchange IP packets. Communication with other machines generally requires further modification to routing tables and/or ARP (Address Resolution Protocol) tables. In most cases the *defaultroute* and/or *proxyarp* options are sufficient for this, but in some cases further intervention is required. The */etc/ppp/ip-up* script can be used for this.

Sometimes it is desirable to add a default route through the remote host, as in the case of a machine whose only connection to the Internet is through the ppp interface. The *defaultroute* option causes pppd to create such a default route when IPCP comes up, and delete it when the link is terminated.

In some cases it is desirable to use proxy ARP, for example on a server machine connected to a LAN, in order to allow other hosts to communicate with the remote host. The *proxyarp* option causes pppd to look for

a network interface on the same subnet as the remote host (an interface supporting broadcast and ARP, which is up and not a point-to-point or loopback interface). If found, `pppd` creates a permanent, published ARP entry with the IP address of the remote host and the hardware address of the network interface found.

When the *demand* option is used, the interface IP addresses have already been set at the point when IPCP comes up. If `pppd` has not been able to negotiate the same addresses that it used to configure the interface (for example when the peer is an ISP that uses dynamic IP address assignment), `pppd` has to change the interface IP addresses to the negotiated addresses. This may disrupt existing connections, and the use of demand dialling with peers that do dynamic IP address assignment is not recommended.

MULTILINK

Multilink PPP provides the capability to combine two or more PPP links between a pair of machines into a single 'bundle', which appears as a single virtual PPP link which has the combined bandwidth of the individual links. Currently, multilink PPP is only supported under Linux.

`Pppd` detects that the link it is controlling is connected to the same peer as another link using the peer's endpoint discriminator and the authenticated identity of the peer (if it authenticates itself). The endpoint discriminator is a block of data which is hopefully unique for each peer. Several types of data can be used, including locally-assigned strings of bytes, IP addresses, MAC addresses, randomly strings of bytes, or E-164 phone numbers. The endpoint discriminator sent to the peer by `pppd` can be set using the endpoint option.

In some circumstances the peer may send no endpoint discriminator or a non-unique value. The bundle option adds an extra string which is added to the peer's endpoint discriminator and authenticated identity when matching up links to be joined together in a bundle. The bundle option can also be used to allow the establishment of multiple bundles between the local system and the peer. `Pppd` uses a TDB database in `/var/run/pppd2.tdb` to match up links.

Assuming that multilink is enabled and the peer is willing to negotiate multilink, then when `pppd` is invoked to bring up the first link to the peer, it will detect that no other link is connected to the peer and create a new bundle, that is, another ppp network interface unit. When another `pppd` is invoked to bring up another link to the peer, it will detect the existing bundle and join its link to it.

If the first link terminates (for example, because of a hangup or a received LCP terminate-request) the bundle is not destroyed unless there are no other links remaining in the bundle. Rather than exiting, the first `pppd` keeps running after its link terminates, until all the links in the bundle have terminated. If the first `pppd` receives a SIGTERM or SIGINT signal, it will destroy the bundle and send a SIGHUP to the `pppd` processes for each of the links in the bundle. If the first `pppd` receives a SIGHUP signal, it will terminate its link but not the bundle.

Note: demand mode is not currently supported with multilink.

EXAMPLES

The following examples assume that the `/etc/ppp/options` file contains the *auth* option (as in the default `/etc/ppp/options` file in the ppp distribution).

Probably the most common use of `pppd` is to dial out to an ISP. This can be done with a command such as

```
pppd call isp
```

where the `/etc/ppp/peers/isp` file is set up by the system administrator to contain something like this:

```
ttyS0 19200 crtscts
connect '/usr/sbin/chat -v -f /etc/ppp/chat-isp'
noauth
```

In this example, we are using chat to dial the ISP's modem and go through any logon sequence required. The `/etc/ppp/chat-isp` file contains the script used by chat; it could for example contain something like this:

```
ABORT "NO CARRIER"
ABORT "NO DIALTONE"
ABORT "ERROR"
ABORT "NO ANSWER"
```

```

ABORT "BUSY"
ABORT "Username/Password Incorrect"
"" "at"
OK "at&d0&c1"
OK "atdt2468135"
"name:" "^Umyuserid"
"word:" "\qmypassword"
"ispts" "\q^Uppp"
"~^Uppp~"

```

See the chat(8) man page for details of chat scripts.

Pppd can also be used to provide a dial-in ppp service for users. If the users already have login accounts, the simplest way to set up the ppp service is to let the users log in to their accounts and run pppd (installed setuid-root) with a command such as

```
pppd proxyarp
```

To allow a user to use the PPP facilities, you need to allocate an IP address for that user's machine and create an entry in /etc/ppp/pap-secrets, /etc/ppp/chap-secrets, or /etc/ppp/srp-secrets (depending on which authentication method the PPP implementation on the user's machine supports), so that the user's machine can authenticate itself. For example, if Joe has a machine called "joespc" that is to be allowed to dial in to the machine called "server" and use the IP address joespc.my.net, you would add an entry like this to /etc/ppp/pap-secrets or /etc/ppp/chap-secrets:

```
joespc  server  "joe's secret"  joespc.my.net
```

(See srp-entry(8) for a means to generate the server's entry when SRP-SHA1 is in use.) Alternatively, you can create a username called (for example) "ppp", whose login shell is pppd and whose home directory is /etc/ppp. Options to be used when pppd is run this way can be put in /etc/ppp/ppprc.

If your serial connection is any more complicated than a piece of wire, you may need to arrange for some control characters to be escaped. In particular, it is often useful to escape XON (^Q) and XOFF (^S), using *asyncmap a0000*. If the path includes a telnet, you probably should escape ^] as well (*asyncmap 200a0000*). If the path includes an rlogin, you will need to use the *escape ff* option on the end which is running the rlogin client, since many rlogin implementations are not transparent; they will remove the sequence [0xff, 0xff, 0x73, 0x73, followed by any 8 bytes] from the stream.

DIAGNOSTICS

Messages are sent to the syslog daemon using facility LOG_DAEMON. (This can be overridden by recompiling pppd with the macro LOG_PPP defined as the desired facility.) See the syslog(8) documentation for details of where the syslog daemon will write the messages. On most systems, the syslog daemon uses the /etc/syslog.conf file to specify the destination(s) for syslog messages. You may need to edit that file to suit.

The *debug* option causes the contents of all control packets sent or received to be logged, that is, all LCP, PAP, CHAP, EAP, or IPCP packets. This can be useful if the PPP negotiation does not succeed or if authentication fails. If debugging is enabled at compile time, the *debug* option also causes other debugging messages to be logged.

Debugging can also be enabled or disabled by sending a SIGUSR1 signal to the pppd process. This signal acts as a toggle.

EXIT STATUS

The exit status of pppd is set to indicate whether any error was detected, or the reason for the link being terminated. The values used are:

- 0** Pppd has detached, or otherwise the connection was successfully established and terminated at the peer's request.
- 1** An immediately fatal error of some kind occurred, such as an essential system call failing, or running out of virtual memory.

- 2 An error was detected in processing the options given, such as two mutually exclusive options being used.
- 3 Pppd is not setuid-root and the invoking user is not root.
- 4 The kernel does not support PPP, for example, the PPP kernel driver is not included or cannot be loaded.
- 5 Pppd terminated because it was sent a SIGINT, SIGTERM or SIGHUP signal.
- 6 The serial port could not be locked.
- 7 The serial port could not be opened.
- 8 The connect script failed (returned a non-zero exit status).
- 9 The command specified as the argument to the *pty* option could not be run.
- 10 The PPP negotiation failed, that is, it didn't reach the point where at least one network protocol (e.g. IP) was running.
- 11 The peer system failed (or refused) to authenticate itself.
- 12 The link was established successfully and terminated because it was idle.
- 13 The link was established successfully and terminated because the connect time limit was reached.
- 14 Callback was negotiated and an incoming call should arrive shortly.
- 15 The link was terminated because the peer is not responding to echo requests.
- 16 The link was terminated by the modem hanging up.
- 17 The PPP negotiation failed because serial loopback was detected.
- 18 The init script failed (returned a non-zero exit status).
- 19 We failed to authenticate ourselves to the peer.

SCRIPTS

Pppd invokes scripts at various stages in its processing which can be used to perform site-specific ancillary processing. These scripts are usually shell scripts, but could be executable code files instead. Pppd does not wait for the scripts to finish (except for the ip-pre-up script). The scripts are executed as root (with the real and effective user-id set to 0), so that they can do things such as update routing tables or run privileged daemons. Be careful that the contents of these scripts do not compromise your system's security. Pppd runs the scripts with standard input, output and error redirected to /dev/null, and with an environment that is empty except for some environment variables that give information about the link. The environment variables that pppd sets are:

DEVICE

The name of the serial tty device being used.

IFNAME

The name of the network interface being used.

IPLOCAL

The IP address for the local end of the link. This is only set when IPCP has come up.

IPREMOTE

The IP address for the remote end of the link. This is only set when IPCP has come up.

PEERNAME

The authenticated name of the peer. This is only set if the peer authenticates itself.

SPEED

The baud rate of the tty device.

ORIG_UID

The real user-id of the user who invoked `pppd`.

PPPLOGNAME

The username of the real user-id that invoked `pppd`. This is always set.

For the `ip-down` and `auth-down` scripts, `pppd` also sets the following variables giving statistics for the connection:

CONNECT_TIME

The number of seconds from when the PPP negotiation started until the connection was terminated.

BYTES_SENT

The number of bytes sent (at the level of the serial port) during the connection.

BYTES_RCVD

The number of bytes received (at the level of the serial port) during the connection.

LINKNAME

The logical name of the link, set with the *linkname* option.

CALL_FILE

The value of the *call* option.

DNS1 If the peer supplies DNS server addresses, this variable is set to the first DNS server address supplied (whether or not the `usepeerdns` option was given).

DNS2 If the peer supplies DNS server addresses, this variable is set to the second DNS server address supplied (whether or not the `usepeerdns` option was given).

`Pppd` invokes the following scripts, if they exist. It is not an error if they don't exist.

/etc/ppp/auth-up

A program or script which is executed after the remote system successfully authenticates itself. It is executed with the parameters

interface-name peer-name user-name tty-device speed

Note that this script is not executed if the peer doesn't authenticate itself, for example when the *noauth* option is used.

/etc/ppp/auth-down

A program or script which is executed when the link goes down, if `/etc/ppp/auth-up` was previously executed. It is executed in the same manner with the same parameters as `/etc/ppp/auth-up`.

/etc/ppp/ip-pre-up

A program or script which is executed just before the ppp network interface is brought up. It is executed with the same parameters as the `ip-up` script (below). At this point the interface exists and has IP addresses assigned but is still down. This can be used to add firewall rules before any IP traffic can pass through the interface. `Pppd` will wait for this script to finish before bringing the interface up, so this script should run quickly.

/etc/ppp/ip-up

A program or script which is executed when the link is available for sending and receiving IP packets (that is, IPCP has come up). It is executed with the parameters

interface-name tty-device speed local-IP-address remote-IP-address ipparam

/etc/ppp/ip-down

A program or script which is executed when the link is no longer available for sending and receiving IP packets. This script can be used for undoing the effects of the `/etc/ppp/ip-up` and `/etc/ppp/ip-pre-up` scripts. It is invoked in the same manner and with the same parameters as the `ip-up` script.

/etc/ppp/ipv6-up

Like /etc/ppp/ip-up, except that it is executed when the link is available for sending and receiving IPv6 packets. It is executed with the parameters

interface-name tty-device speed local-link-local-address remote-link-local-address ipparam

/etc/ppp/ipv6-down

Similar to /etc/ppp/ip-down, but it is executed when IPv6 packets can no longer be transmitted on the link. It is executed with the same parameters as the ipv6-up script.

/etc/ppp/ipx-up

A program or script which is executed when the link is available for sending and receiving IPX packets (that is, IPXCP has come up). It is executed with the parameters

interface-name tty-device speed network-number local-IPX-node-address remote-IPX-node-address local-IPX-routing-protocol remote-IPX-routing-protocol local-IPX-router-name remote-IPX-router-name ipparam pppd-pid

The local-IPX-routing-protocol and remote-IPX-routing-protocol field may be one of the following:

NONE to indicate that there is no routing protocol
 RIP to indicate that RIP/SAP should be used
 NLSP to indicate that Novell NLSP should be used
 RIP NLSP to indicate that both RIP/SAP and NLSP should be used

/etc/ppp/ipx-down

A program or script which is executed when the link is no longer available for sending and receiving IPX packets. This script can be used for undoing the effects of the /etc/ppp/ipx-up script. It is invoked in the same manner and with the same parameters as the ipx-up script.

FILES

/var/run/pppdn.pid (BSD or Linux), **/etc/ppp/pppdn.pid** (others)

Process-ID for pppd process on ppp interface unit *n*.

/var/run/pppd-name.pid (BSD or Linux),

/etc/ppp/pppd-name.pid (others) Process-ID for pppd process for logical link *name* (see the *linkname* option).

/var/run/pppd2.tdb

Database containing information about pppd processes, interfaces and links, used for matching links to bundles in multilink operation. May be examined by external programs to obtain information about running pppd instances, the interfaces and devices they are using, IP address assignments, etc. **/etc/ppp/pap-secrets** Usernames, passwords and IP addresses for PAP authentication. This file should be owned by root and not readable or writable by any other user. Pppd will log a warning if this is not the case.

/etc/ppp/chap-secrets

Names, secrets and IP addresses for CHAP/MS-CHAP/MS-CHAPv2 authentication. As for /etc/ppp/pap-secrets, this file should be owned by root and not readable or writable by any other user. Pppd will log a warning if this is not the case.

/etc/ppp/srp-secrets

Names, secrets, and IP addresses for EAP authentication. As for /etc/ppp/pap-secrets, this file should be owned by root and not readable or writable by any other user. Pppd will log a warning if this is not the case.

~/ppp_pseudonym

Saved client-side SRP-SHA1 pseudonym. See *thesrp-use-pseudonym* option for details.

/etc/ppp/options

System default options for pppd, read before user default options or command-line options.

~/.ppprc

User default options, read before /etc/ppp/options.*ttname*.

/etc/ppp/options.*ttname*

System default options for the serial port being used, read after ~/.ppprc. In forming the *ttname* part of this filename, an initial /dev/ is stripped from the port name (if present), and any slashes in the remaining part are converted to dots.

/etc/ppp/peers

A directory containing options files which may contain privileged options, even if pppd was invoked by a user other than root. The system administrator can create options files in this directory to permit non-privileged users to dial out without requiring the peer to authenticate, but only to certain trusted peers.

SEE ALSO

chat(8), pppstats(8)

RFC1144

Jacobson, V. *Compressing TCP/IP headers for low-speed serial links*. February 1990.

RFC1321

Rivest, R. *The MD5 Message-Digest Algorithm*. April 1992.

RFC1332

McGregor, G. *PPP Internet Protocol Control Protocol (IPCP)*. May 1992.

RFC1334

Lloyd, B.; Simpson, W.A. *PPP authentication protocols*. October 1992.

RFC1661

Simpson, W.A. *The Point-to-Point Protocol (PPP)*. July 1994.

RFC1662

Simpson, W.A. *PPP in HDLC-like Framing*. July 1994.

RFC1990

Sklower, K.; et al., *The PPP Multilink Protocol (MP)*. August 1996.

RFC2284

Blunk, L.; Vollbrecht, J., *PPP Extensible Authentication Protocol (EAP)*. March 1998.

RFC2472

Haskin, D. *IP Version 6 over PPP* December 1998.

RFC2945

Wu, T., *The SRP Authentication and Key Exchange System* September 2000.

draft-ietf-pppext-eap-srp-03.txt

Carlson, J.; et al., *EAP SRP-SHA1 Authentication Protocol*. July 2001.

NOTES

Some limited degree of control can be exercised over a running pppd process by sending it a signal from the list below.

SIGINT, SIGTERM

These signals cause pppd to terminate the link (by closing LCP), restore the serial device settings, and exit. If a connector or disconnector process is currently running, pppd will send the same signal to its process group, so as to terminate the connector or disconnector process.

SIGHUP

This signal causes pppd to terminate the link, restore the serial device settings, and close the serial device. If the *persist* or *demand* option has been specified, pppd will try to reopen the serial device and start another connection (after the holdoff period). Otherwise pppd will exit. If this signal is received during the holdoff period, it causes pppd to end the holdoff period immediately. If a connector or disconnector process is running, pppd will send the same signal to its process group.

SIGUSR1

This signal toggles the state of the *debug* option.

SIGUSR2

This signal causes pppd to renegotiate compression. This can be useful to re-enable compression after it has been disabled as a result of a fatal decompression error. (Fatal decompression errors generally indicate a bug in one or other implementation.)

AUTHORS

Paul Mackerras (paulus@samba.org), based on earlier work by Drew Perkins, Brad Clements, Karl Fox, Greg Christy, and Brad Parker.

COPYRIGHT

Pppd is copyrighted and made available under conditions which provide that it may be copied and used in source or binary forms provided that the conditions listed below are met. Portions of pppd are covered by the following copyright notices:

Copyright (c) 1984-2000 Carnegie Mellon University. All rights reserved.

Copyright (c) 1993-2004 Paul Mackerras. All rights reserved.

Copyright (c) 1995 Pedro Roque Marques. All rights reserved.

Copyright (c) 1995 Eric Rosenquist. All rights reserved.

Copyright (c) 1999 Tommi Komulainen. All rights reserved.

Copyright (C) Andrew Tridgell 1999

Copyright (c) 2000 by Sun Microsystems, Inc. All rights reserved.

Copyright (c) 2001 by Sun Microsystems, Inc. All rights reserved.

Copyright (c) 2002 Google, Inc. All rights reserved.

The copyright notices contain the following statements.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
3. The name "Carnegie Mellon University" must not be used to endorse or promote products derived from this software without prior written permission. For permission or any legal details, please contact
Office of Technology Transfer
Carnegie Mellon University
5000 Forbes Avenue
Pittsburgh, PA 15213-3890
(412) 268-4387, fax: (412) 268-7395
tech-transfer@andrew.cmu.edu
- 3b. The name(s) of the authors of this software must not be used to endorse or promote products derived from this software without prior written permission.
4. Redistributions of any form whatsoever must retain the following acknowledgements:
"This product includes software developed by Computing Services at Carnegie Mellon University (<http://www.cmu.edu/computing/>)."

"This product includes software developed by Paul Mackerras
<paulus@samba.org>".

"This product includes software developed by Pedro Roque Marques
<pedro_m@yahoo.com>".

"This product includes software developed by Tommi Komulainen
<Tommi.Komulainen@iki.fi>".

CARNEGIE MELLON UNIVERSITY DISCLAIMS ALL WARRANTIES WITH REGARD TO THIS SOFTWARE, INCLUDING ALL IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS, IN NO EVENT SHALL CARNEGIE MELLON UNIVERSITY BE LIABLE FOR ANY SPECIAL, INDIRECT OR CONSEQUENTIAL DAMAGES OR ANY DAMAGES WHATSOEVER RESULTING FROM LOSS OF USE, DATA OR PROFITS, WHETHER IN AN ACTION OF CONTRACT, NEGLIGENCE OR OTHER TORTIOUS ACTION, ARISING OUT OF OR IN CONNECTION WITH THE USE OR PERFORMANCE OF THIS SOFTWARE.

THE AUTHORS OF THIS SOFTWARE DISCLAIM ALL WARRANTIES WITH REGARD TO THIS SOFTWARE, INCLUDING ALL IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS, IN NO EVENT SHALL THE AUTHORS BE LIABLE FOR ANY SPECIAL, INDIRECT OR CONSEQUENTIAL DAMAGES OR ANY DAMAGES WHATSOEVER RESULTING FROM LOSS OF USE, DATA OR PROFITS, WHETHER IN AN ACTION OF CONTRACT, NEGLIGENCE OR OTHER TORTIOUS ACTION, ARISING OUT OF OR IN CONNECTION WITH THE USE OR PERFORMANCE OF THIS SOFTWARE.