

NAME

`selinux_status_open`, `selinux_status_close`, `selinux_status_updated`, `selinux_status_getenforce`, `selinux_status_policyload` and `selinux_status_deny_unknown` – reference the SELinux kernel status without invocation of system calls

SYNOPSIS

```
#include <selinux/avc.h>
```

```
int selinux_status_open(int fallback);
```

```
void selinux_status_close(void);
```

```
int selinux_status_updated(void);
```

```
int selinux_status_getenforce(void);
```

```
int selinux_status_policyload(void);
```

```
int selinux_status_deny_unknown(void);
```

DESCRIPTION

Linux 2.6.37 or later provides a SELinux kernel status page; being mostly placed on `/sys/fs/selinux/status` entry. It enables userspace applications to mmap this page with read-only mode, then it informs some status without system call invocations.

In some cases that a userspace application tries to apply heavy frequent access control; such as row-level security in databases, it will face unignorable cost to communicate with kernel space to check invalidation of userspace avc.

These functions provides applications a way to know some kernel events without system-call invocation or worker thread for monitoring.

`selinux_status_open()` tries to **`open(2)`** `/sys/fs/selinux/status` and **`mmap(2)`** it in read-only mode. The file-descriptor and pointer to the page shall be stored internally; Don't touch them directly. Set 1 on the *fallback* argument to handle a case of older kernels without kernel status page support. In this case, this function tries to open a netlink socket using **`avc_netlink_open(3)`** and overwrite corresponding callbacks (`setenforce` and `policyload`). Thus, we need to pay attention to the interaction with these interfaces, when fallback mode is enabled.

`selinux_status_close()` unmap the kernel status page and close its file descriptor, or close the netlink socket if fallbacked.

`selinux_status_updated()` processes status update events. There are two kinds of status updates. **`setenforce`** events will change the effective enforcing state used within the AVC, and **`policyload`** events will result in a cache flush.

This function returns 0 if there have been no updates since the last call, 1 if there have been updates since the last call, or -1 on error.

`selinux_status_getenforce()` returns 0 if SELinux is running in permissive mode, 1 if enforcing mode, or -1 on error. Same as **`assecurity_getenforce(3)`** except with or without system call invocation.

`selinux_status_policyload()` returns times of policy reloaded on the running system, or -1 on error. Note that it is not a reliable value on fallback-mode until it receive the first event message via netlink socket. Thus, don't use this value to know actual times of policy reloaded.

selinux_status_deny_unknown() returns 0 if SELinux treats policy queries on undefined object classes or permissions as being allowed, 1 if such queries are denied, or -1 on error.

Also note that these interfaces are not thread-safe, so you have to protect them from concurrent calls using exclusive locks when multiple threads are performing.

RETURN VALUE

selinux_status_open() returns 0 or 1 on success. 1 means we are ready to use these interfaces, but netlink socket was opened as fallback instead of the kernel status page. On error, -1 shall be returned.

Any other functions with a return value shall return its characteristic value as described above, or -1 on errors.

SEE ALSO

mmap(2), **avc_netlink_open(3)**, **security_getenforce(3)**, **security_deny_unknown(3)**