

NAME

get_thread_area, set_thread_area – manipulate thread-local storage information

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <sys/syscall.h> /* Definition of SYS_* constants */
#include <unistd.h>

#ifdef __i386__ || defined __x86_64__
# include <asm/ldt.h> /* Definition of struct user_desc */

int syscall(SYS_get_thread_area, struct user_desc *u_info);
int syscall(SYS_set_thread_area, struct user_desc *u_info);

#elif defined __m68k__

int syscall(SYS_get_thread_area);
int syscall(SYS_set_thread_area, unsigned long tp);

#elif defined __mips__

int syscall(SYS_set_thread_area, unsigned long addr);

#endif
```

Note: glibc provides no wrappers for these system calls, necessitating the use of **syscall(2)**.

DESCRIPTION

These calls provide architecture-specific support for a thread-local storage implementation. At the moment, **set_thread_area()** is available on m68k, MIPS, and x86 (both 32-bit and 64-bit variants); **get_thread_area()** is available on m68k and x86.

On m68k and MIPS, **set_thread_area()** allows storing an arbitrary pointer (provided in the **tp** argument on m68k and in the **addr** argument on MIPS) in the kernel data structure associated with the calling thread; this pointer can later be retrieved using **get_thread_area()** (see also NOTES for information regarding obtaining the thread pointer on MIPS).

On x86, Linux dedicates three global descriptor table (GDT) entries for thread-local storage. For more information about the GDT, see the Intel Software Developer's Manual or the AMD Architecture Programming Manual.

Both of these system calls take an argument that is a pointer to a structure of the following type:

```
struct user_desc {
    unsigned int  entry_number;
    unsigned int  base_addr;
    unsigned int  limit;
    unsigned int  seg_32bit:1;
    unsigned int  contents:2;
    unsigned int  read_exec_only:1;
    unsigned int  limit_in_pages:1;
    unsigned int  seg_not_present:1;
    unsigned int  useable:1;
#ifdef __x86_64__
    unsigned int  lm:1;
#endif
};
```

get_thread_area() reads the GDT entry indicated by *u_info->entry_number* and fills in the rest of the fields in *u_info*.

set_thread_area() sets a TLS entry in the GDT.

The TLS array entry set by **set_thread_area()** corresponds to the value of *u_info->entry_number* passed in by the user. If this value is in bounds, **set_thread_area()** writes the TLS descriptor pointed to by *u_info* into the thread's TLS array.

When **set_thread_area()** is passed an *entry_number* of `-1`, it searches for a free TLS entry. If **set_thread_area()** finds a free TLS entry, the value of *u_info->entry_number* is set upon return to show which entry was changed.

A *user_desc* is considered "empty" if *read_exec_only* and *seg_not_present* are set to 1 and all of the other fields are 0. If an "empty" descriptor is passed to **set_thread_area()**, the corresponding TLS entry will be cleared. See BUGS for additional details.

Since Linux 3.19, **set_thread_area()** cannot be used to write non-present segments, 16-bit segments, or code segments, although clearing a segment is still acceptable.

RETURN VALUE

On x86, these system calls return 0 on success, and `-1` on failure, with *errno* set to indicate the error.

On MIPS and m68k, **set_thread_area()** always returns 0. On m68k, **get_thread_area()** returns the thread area pointer value (previously set via **set_thread_area()**).

ERRORS

EFAULT

u_info is an invalid pointer.

EINVAL

u_info->entry_number is out of bounds.

ENOSYS

get_thread_area() or **set_thread_area()** was invoked as a 64-bit system call.

ESRCH

(**set_thread_area()**) A free TLS entry could not be located.

VERSIONS

set_thread_area() first appeared in Linux 2.5.29. **get_thread_area()** first appeared in Linux 2.5.32.

STANDARDS

set_thread_area() and **get_thread_area()** are Linux-specific and should not be used in programs that are intended to be portable.

NOTES

These system calls are generally intended for use only by threading libraries.

arch_prctl(2) can interfere with **set_thread_area()** on x86. See **arch_prctl(2)** for more details. This is not normally a problem, as **arch_prctl(2)** is normally used only by 64-bit programs.

On MIPS, the current value of the thread area pointer can be obtained using the instruction:

```
rdhwr dest, $29
```

This instruction traps and is handled by kernel.

BUGS

On 64-bit kernels before Linux 3.19, one of the padding bits in *user_desc*, if set, would prevent the descriptor from being considered empty (see **modify_ldt(2)**). As a result, the only reliable way to clear a TLS entry is to use **memset(3)** to zero the entire *user_desc* structure, including padding bits, and then to set the *read_exec_only* and *seg_not_present* bits. On Linux 3.19, a *user_desc* consisting entirely of zeros except for *entry_number* will also be interpreted as a request to clear a TLS entry, but this behaved differently on older kernels.

Prior to Linux 3.19, the DS and ES segment registers must not reference TLS entries.

SEE ALSO

arch_prctl(2), **modify_ldt(2)**, **ptrace(2)** (**PTRACE_GET_THREAD_AREA** and **PTRACE_SET_THREAD_AREA**)