

NAME

getpriority, setpriority – get/set program scheduling priority

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <sys/resource.h>
```

```
int getpriority(int which, id_t who);
```

```
int setpriority(int which, id_t who, int prio);
```

DESCRIPTION

The scheduling priority of the process, process group, or user, as indicated by *which* and *who* is obtained with the **getpriority()** call and set with the **setpriority()** call. The process attribute dealt with by these system calls is the same attribute (also known as the "nice" value) that is dealt with by **nice(2)**.

The value *which* is one of **PRIO_PROCESS**, **PRIO_PGRP**, or **PRIO_USER**, and *who* is interpreted relative to *which* (a process identifier for **PRIO_PROCESS**, process group identifier for **PRIO_PGRP**, and a user ID for **PRIO_USER**). A zero value for *who* denotes (respectively) the calling process, the process group of the calling process, or the real user ID of the calling process.

The *prio* argument is a value in the range -20 to 19 (but see NOTES below), with -20 being the highest priority and 19 being the lowest priority. Attempts to set a priority outside this range are silently clamped to the range. The default priority is 0 ; lower values give a process a higher scheduling priority.

The **getpriority()** call returns the highest priority (lowest numerical value) enjoyed by any of the specified processes. The **setpriority()** call sets the priorities of all of the specified processes to the specified value.

Traditionally, only a privileged process could lower the nice value (i.e., set a higher priority). However, since Linux 2.6.12, an unprivileged process can decrease the nice value of a target process that has a suitable **RLIMIT_NICE** soft limit; see **getrlimit(2)** for details.

RETURN VALUE

On success, **getpriority()** returns the calling thread's nice value, which may be a negative number. On error, it returns -1 and sets *errno* to indicate the error.

Since a successful call to **getpriority()** can legitimately return the value -1 , it is necessary to clear *errno* prior to the call, then check *errno* afterward to determine if -1 is an error or a legitimate value.

setpriority() returns 0 on success. On failure, it returns -1 and sets *errno* to indicate the error.

ERRORS**EACCES**

The caller attempted to set a lower nice value (i.e., a higher process priority), but did not have the required privilege (on Linux: did not have the **CAP_SYS_NICE** capability).

EINVAL

which was not one of **PRIO_PROCESS**, **PRIO_PGRP**, or **PRIO_USER**.

EPERM

A process was located, but its effective user ID did not match either the effective or the real user ID of the caller, and was not privileged (on Linux: did not have the **CAP_SYS_NICE** capability). But see NOTES below.

ESRCH

No process was located using the *which* and *who* values specified.

STANDARDS

POSIX.1-2001, POSIX.1-2008, SVr4, 4.4BSD (these interfaces first appeared in 4.2BSD).

NOTES

For further details on the nice value, see **sched(7)**.

Note: the addition of the "autogroup" feature in Linux 2.6.38 means that the nice value no longer has its

traditional effect in many circumstances. For details, see **sched(7)**.

A child created by **fork(2)** inherits its parent's nice value. The nice value is preserved across **execve(2)**.

The details on the condition for **EPERM** depend on the system. The above description is what POSIX.1-2001 says, and seems to be followed on all System V-like systems. Linux kernels before Linux 2.6.12 required the real or effective user ID of the caller to match the real user of the process *who* (instead of its effective user ID). Linux 2.6.12 and later require the effective user ID of the caller to match the real or effective user ID of the process *who*. All BSD-like systems (SunOS 4.1.3, Ultrix 4.2, 4.3BSD, FreeBSD 4.3, OpenBSD-2.5, ...) behave in the same manner as Linux 2.6.12 and later.

C library/kernel differences

The getpriority system call returns nice values translated to the range 40..1, since a negative return value would be interpreted as an error. The glibc wrapper function **forgetpriority()** translates the value back according to the formula $unice = 20 - knice$ (thus, the 40..1 range returned by the kernel corresponds to the range -20..19 as seen by user space).

BUGS

According to POSIX, the nice value is a per-process setting. However, under the current Linux/NPTL implementation of POSIX threads, the nice value is a per-thread attribute: different threads in the same process can have different nice values. Portable applications should avoid relying on the Linux behavior, which may be made standards conformant in the future.

SEE ALSO

nice(1), **renice(1)**, **fork(2)**, **capabilities(7)**, **sched(7)**

Documentation/scheduler/sched-nice-design.txt in the Linux kernel source tree (since Linux 2.6.23)