NAME
chroot – change root directory

LIBRARY
Standard C library (*libc*, *−lc*)

SYNOPSIS
#include <unistd.h>

int chroot(const char **path***);**

Feature Test Macro Requirements for glibc (see **feature_test_macros**(7)):

chroot():
    Since glibc 2.2.2:
        _XOPEN_SOURCE && ! (_POSIX_C_SOURCE >= 200112L)
          || /* Since glibc 2.20: */ _DEFAULT_SOURCE
          || /* glibc <= 2.19: */ _BSD_SOURCE
    Before glibc 2.2.2:
        none

DESCRIPTION
chroot() changes the root directory of the calling process to that specified in *path*. This directory will be used for pathnames beginning with */*. The root directory is inherited by all children of the calling process.

Only a privileged process (Linux: one with the **CAP_SYS_CHROOT** capability in its user namespace) may call **chroot**().

This call changes an ingredient in the pathname resolution process and does nothing else. In particular, it is not intended to be used for any kind of security purpose, neither to fully sandbox a process nor to restrict filesystem system calls. In the past, **chroot**() has been used by daemons to restrict themselves prior to passing paths supplied by untrusted users to system calls such as **open**(2). However, if a folder is moved out of the chroot directory, an attacker can exploit that to get out of the chroot directory as well. The easiest way to do that is to **chdir**(2) to the to-be-moved directory, wait for it to be moved out, then open a path like ../../../etc/passwd.

A slightly trickier variation also works under some circumstances if **chdir**(2) is not permitted. If a daemon allows a "chroot directory" to be specified, that usually means that if you want to prevent remote users from accessing files outside the chroot directory, you must ensure that folders are never moved out of it.

This call does not change the current working directory, so that after the call '.' can be outside the tree rooted at '/'. In particular, the superuser can escape from a "chroot jail" by doing:

```
mkdir foo; chroot foo; cd ..
```

This call does not close open file descriptors, and such file descriptors may allow access to files outside the chroot tree.

RETURN VALUE
On success, zero is returned. On error, −1 is returned, and *errno* is set to indicate the error.

ERRORS
Depending on the filesystem, other errors can be returned. The more general errors are listed below:

EACCES
    Search permission is denied on a component of the path prefix. (See also **path_resolution**(7).)

EFAULT
    path points outside your accessible address space.

EIO   An I/O error occurred.

ELOOP
    Too many symbolic links were encountered in resolving *path*.

ENAMETOOLONG
   path is too long.

ENOENT
   The file does not exist.

ENOMEM
   Insufficient kernel memory was available.

ENOTDIR
   A component of *path* is not a directory.

EPERM
   The caller has insufficient privilege.

STANDARDS
SVr4, 4.4BSD, SUSv2 (marked LEGACY). This function is not part of POSIX.1-2001.

NOTES
A child process created via **fork**(2) inherits its parent's root directory. The root directory is left unchanged by **execve**(2).

The magic symbolic link, */proc/[pid]/root*, can be used to discover a process's root directory; see **proc**(5) for details.

FreeBSD has a stronger **jail**() system call.

SEE ALSO
chroot(1), **chdir**(2), **pivot_root**(2), **path_resolution**(7), **switch_root**(8)