

**NAME**

times – get process times

**LIBRARY**

Standard C library (*libc*, *-lc*)

**SYNOPSIS**

```
#include <sys/times.h>
```

```
clock_t times(struct tms *buf);
```

**DESCRIPTION**

**times()** stores the current process times in the *struct tms* that *buf* points to. The *struct tms* is as defined in *<sys/times.h>*:

```
struct tms {
    clock_t tms_utime; /* user time */
    clock_t tms_stime; /* system time */
    clock_t tms_cutime; /* user time of children */
    clock_t tms_cstime; /* system time of children */
};
```

The *tms\_utime* field contains the CPU time spent executing instructions of the calling process. The *tms\_stime* field contains the CPU time spent executing inside the kernel while performing tasks on behalf of the calling process.

The *tms\_cutime* field contains the sum of the *tms\_utime* and *tms\_cutime* values for all waited-for terminated children. The *tms\_cstime* field contains the sum of the *tms\_stime* and *tms\_cstime* values for all waited-for terminated children.

Times for terminated children (and their descendants) are added in at the moment **wait(2)** or **waitpid(2)** returns their process ID. In particular, times of grandchildren that the children did not wait for are never seen.

All times reported are in clock ticks.

**RETURN VALUE**

**times()** returns the number of clock ticks that have elapsed since an arbitrary point in the past. The return value may overflow the possible range of type *clock\_t*. On error, *(clock\_t) -1* is returned, and *errno* is set to indicate the error.

**ERRORS****EFAULT**

*tms* points outside the process's address space.

**STANDARDS**

POSIX.1-2001, POSIX.1-2008, SVr4, 4.3BSD.

**NOTES**

The number of clock ticks per second can be obtained using:

```
sysconf(_SC_CLK_TCK);
```

In POSIX.1-1996 the symbol **CLK\_TCK** (defined in *<time.h>*) is mentioned as obsolescent. It is obsolete now.

Before Linux 2.6.9, if the disposition of **SIGCHLD** is set to **SIG\_IGN**, then the times of terminated children are automatically included in the *tms\_cstime* and *tms\_cutime* fields, although POSIX.1-2001 says that this should happen only if the calling process **wait(2)**s on its children. This nonconformance is rectified in Linux 2.6.9 and later.

On Linux, the *buf* argument can be specified as **NULL**, with the result that **times()** just returns a function result. However, POSIX does not specify this behavior, and most other UNIX implementations require a non-NULL value for *buf*.

Note that **clock(3)** also returns a value of type *clock\_t*, but this value is measured in units of

**CLOCKS\_PER\_SEC**, not the clock ticks used by **times()**.

On Linux, the “arbitrary point in the past” from which the return value of **times()** is measured has varied across kernel versions. On Linux 2.4 and earlier, this point is the moment the system was booted. Since Linux 2.6, this point is  $(2^{32}/\text{HZ}) - 300$  seconds before system boot time. This variability across kernel versions (and across UNIX implementations), combined with the fact that the returned value may overflow the range of *clock\_t*, means that a portable application would be wise to avoid using this value. To measure changes in elapsed time, use **clock\_gettime(2)** instead.

#### Historical

SVr1-3 returns *long* and the struct members are of type *time\_t* although they store clock ticks, not seconds since the Epoch. V7 used *long* for the struct members, because it had no type *time\_t* yet.

#### BUGS

A limitation of the Linux system call conventions on some architectures (notably i386) means that on Linux 2.6 there is a small time window (41 seconds) soon after boot when **times()** can return  $-1$ , falsely indicating that an error occurred. The same problem can occur when the return value wraps past the maximum value that can be stored in **clock\_t**.

#### SEE ALSO

**time(1)**, **getrusage(2)**, **wait(2)**, **clock(3)**, **sysconf(3)**, **time(7)**